

# **COMP 3331/9331:**

## **Comp** Assignment Project Exam Help **ks and**

<https://eduassistpro.github.io/>

Add WeChat `edu_assist_pro`  
Wee

## Transport Layer (Continued)

Reading Guide: Chapter 3, Sections: 3.4, 3.5

# Transport Layer Outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

■ segment structure  
able data transfer

https://eduassistpro.github.io/  
Add WeChat edu\_assist\_pro  
3.6 mechanisms of congestion control

3.7 TCP congestion control

# rdt2.0 has a fatal flaw!

what happens if  
ACK/NAK corrupted?

- ❖ sender doesn't know what happened at receiver!
- ❖ can't just retransmit

Assignment Project Exam Help

handling duplicates:

- ❖ sender retransmits current pkt if ACK/NAK corrupted

<https://eduassistpro.github.io/>  
Add WeChat: [edu\\_assist\\_pro](https://edu_assist_pro) (doesn't deliver up) duplicate pkt

stop and wait  
sender sends one packet,  
then waits for receiver  
response

# rdt2.1: discussion

## sender:

- ❖ seq # added to pkt
- ❖ two seq. #  
~~Assignment~~ will suffice. Why?
- ❖ must check if ACK/NAK corrupted
- ❖ twice as many states
  - state must “remember” whether “expected” pkt should have seq # of 0 or 1

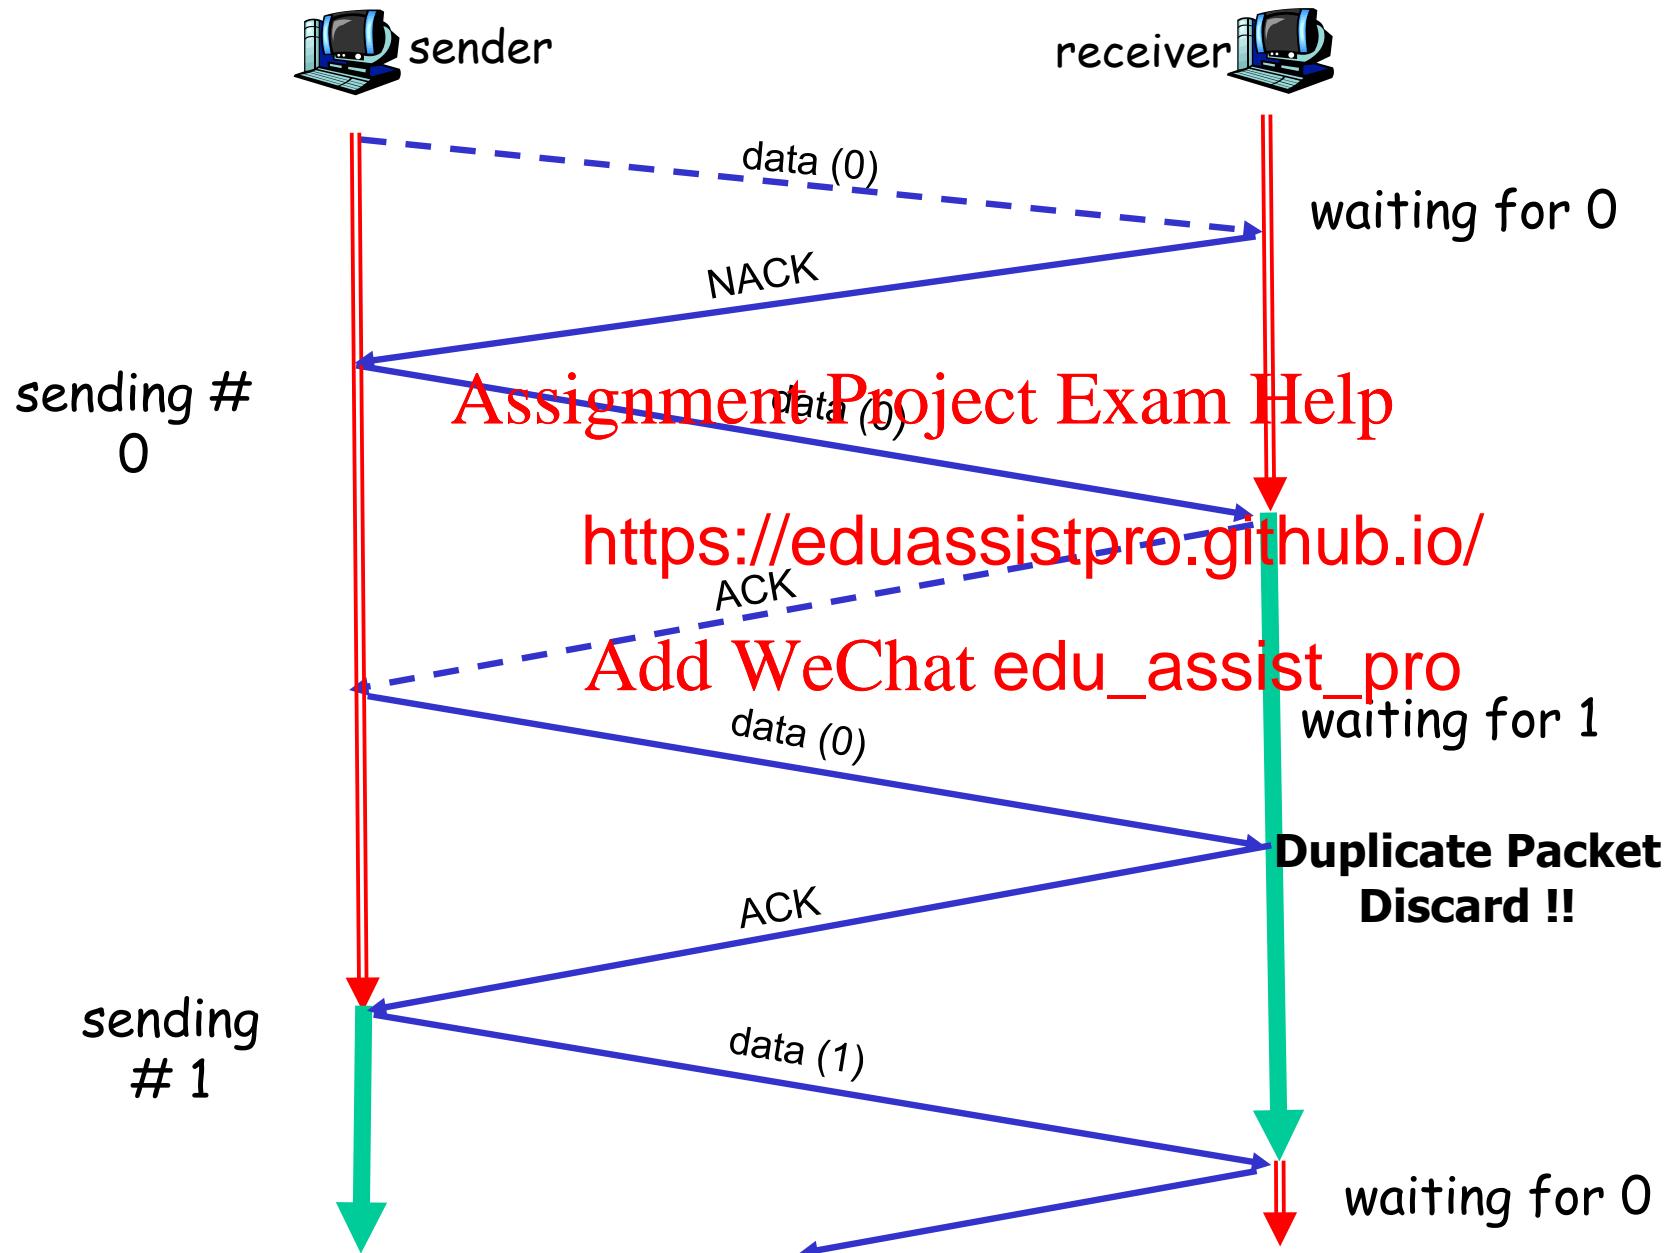
## receiver:

- ❖ must check if received packet is duplicate
- ❖ indicates whether it is expected pkt
- ❖ receiver can not know if its last ACK/NAK received OK at sender

- New Measures: Sequence Numbers, Checksum for ACK/NACK, Duplicate detection

# Another Look at rdt2.1

Dotted line: erroneous transmission  
Solid line: error-free transmission

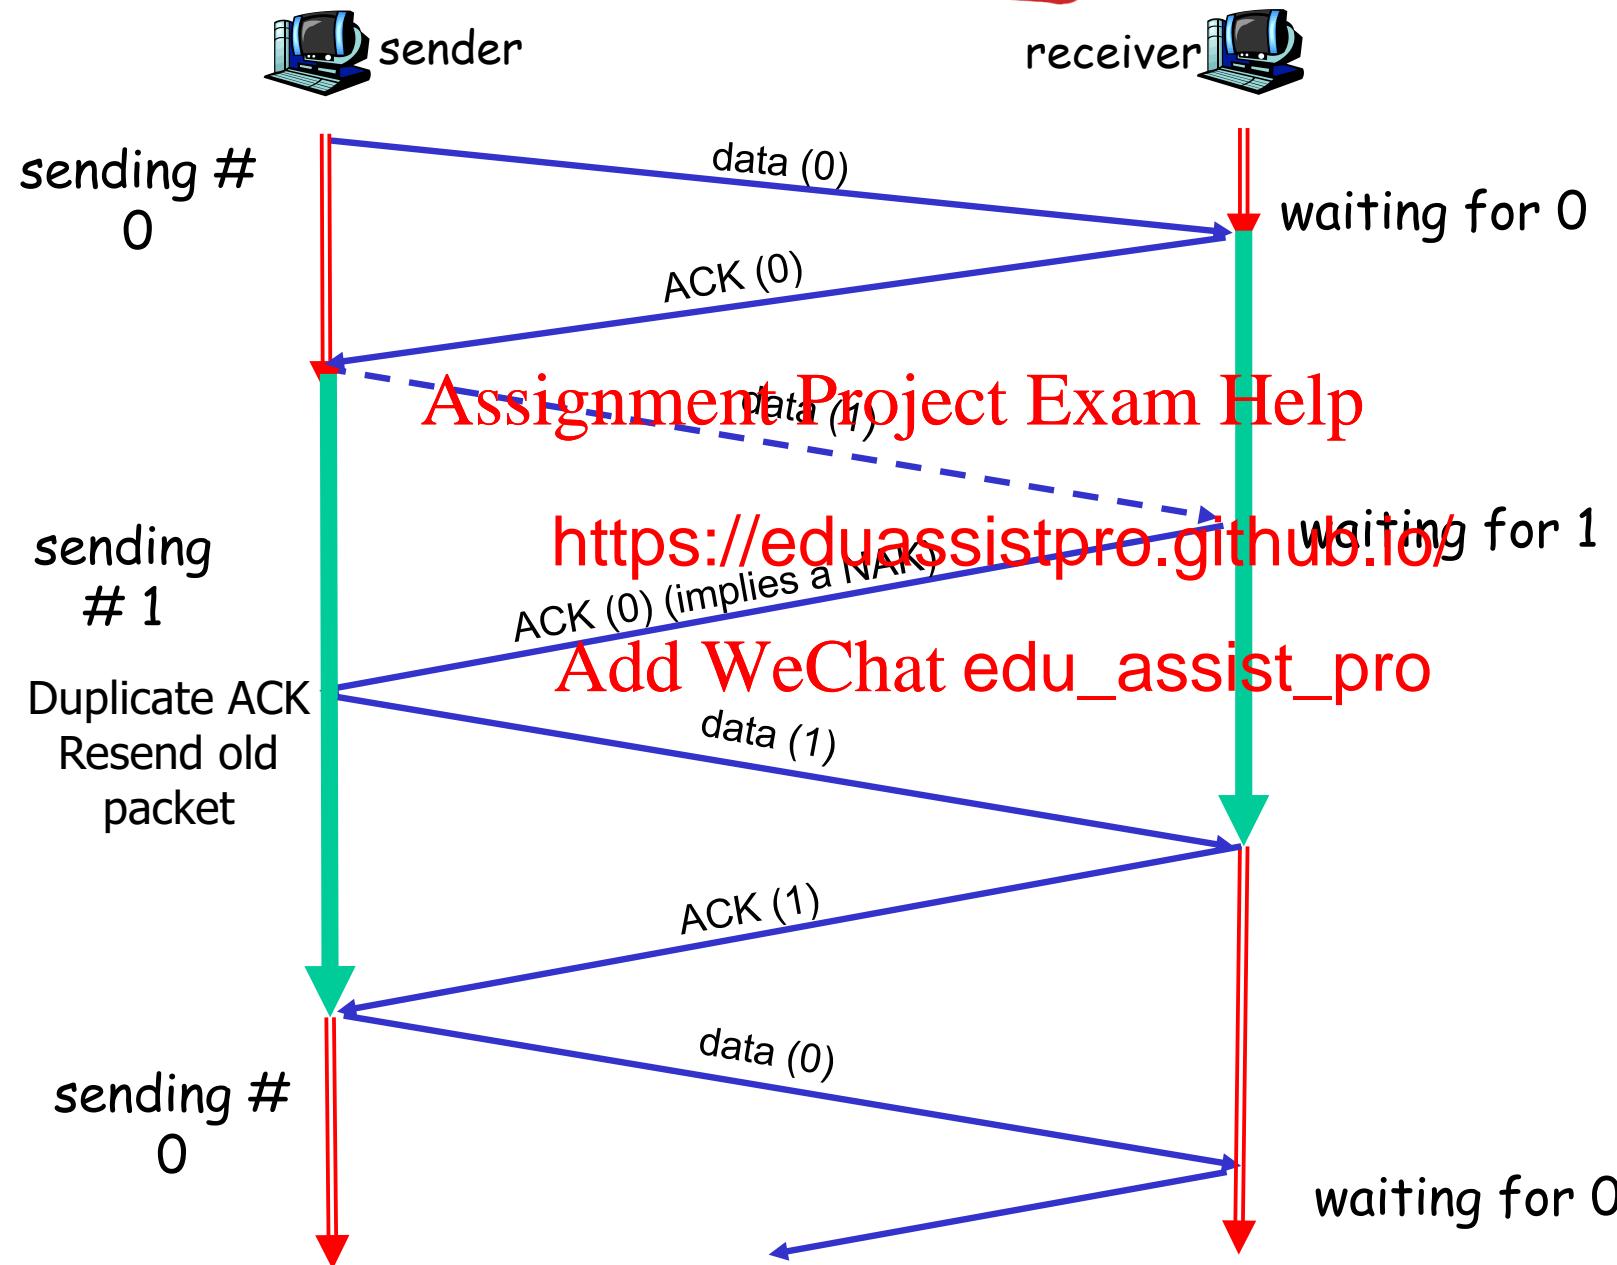


## rdt2.2: a NAK-free protocol

- ❖ same functionality as rdt2.1, using ACKs only
- ❖ instead of NAK, receiver sends ACK for last pkt received OK
  - receiver must f pkt being ACKed
- ❖ duplicate ACK <https://eduassistpro.github.io/>ame action as NAK: *retransmit current pkt*  
[Add WeChat edu\\_assist\\_pro](#)

# rdt2.2: Example

Dotted line: erroneous transmission  
Solid line: error-free transmission



# rdt3.0: channels with errors and loss

## new assumption:

underlying channel can  
also lose packets  
(data, ACKs)

- checksum, see <https://eduassistpro.github.io/>  
ACKs, retransmission will be of help... but not enough

## approach: sender waits

“reasonable” amount of time for ACK  
❖ retransmits if no ACK

in this time  
r ACK

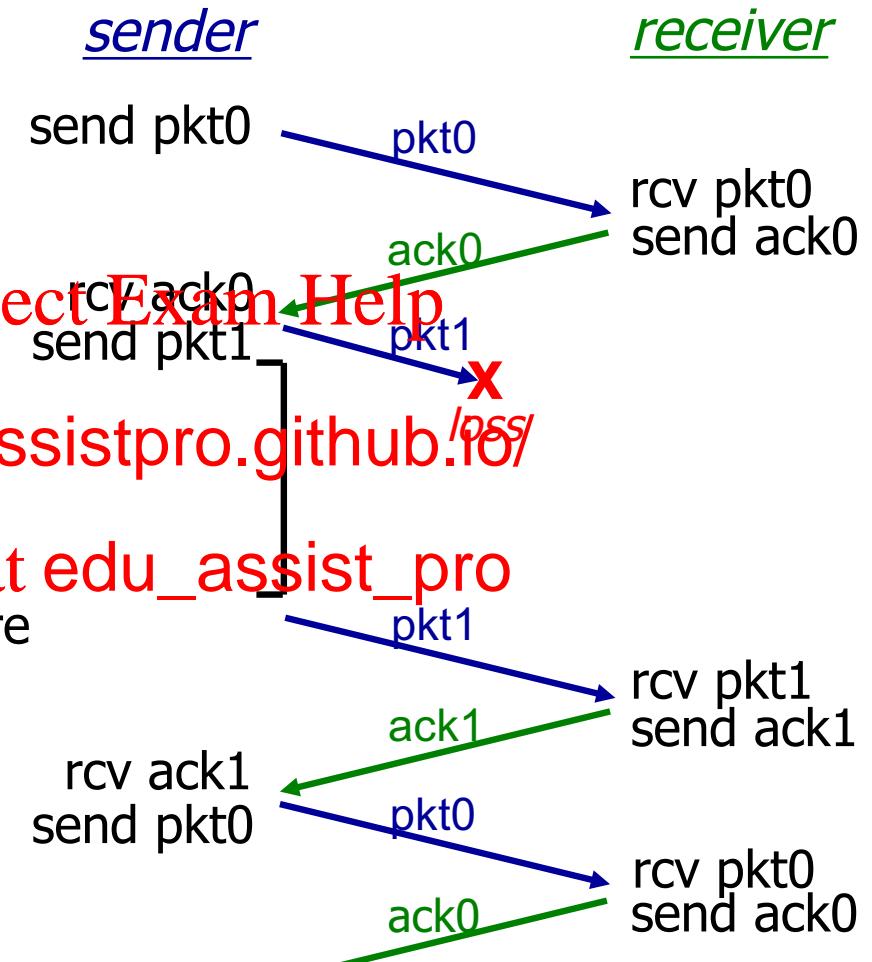
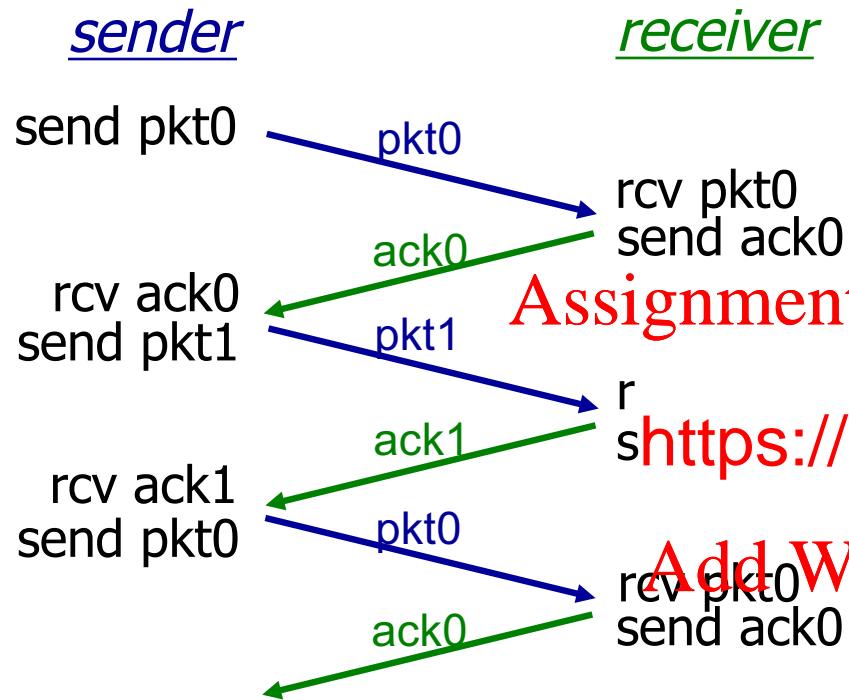
d

mission will be

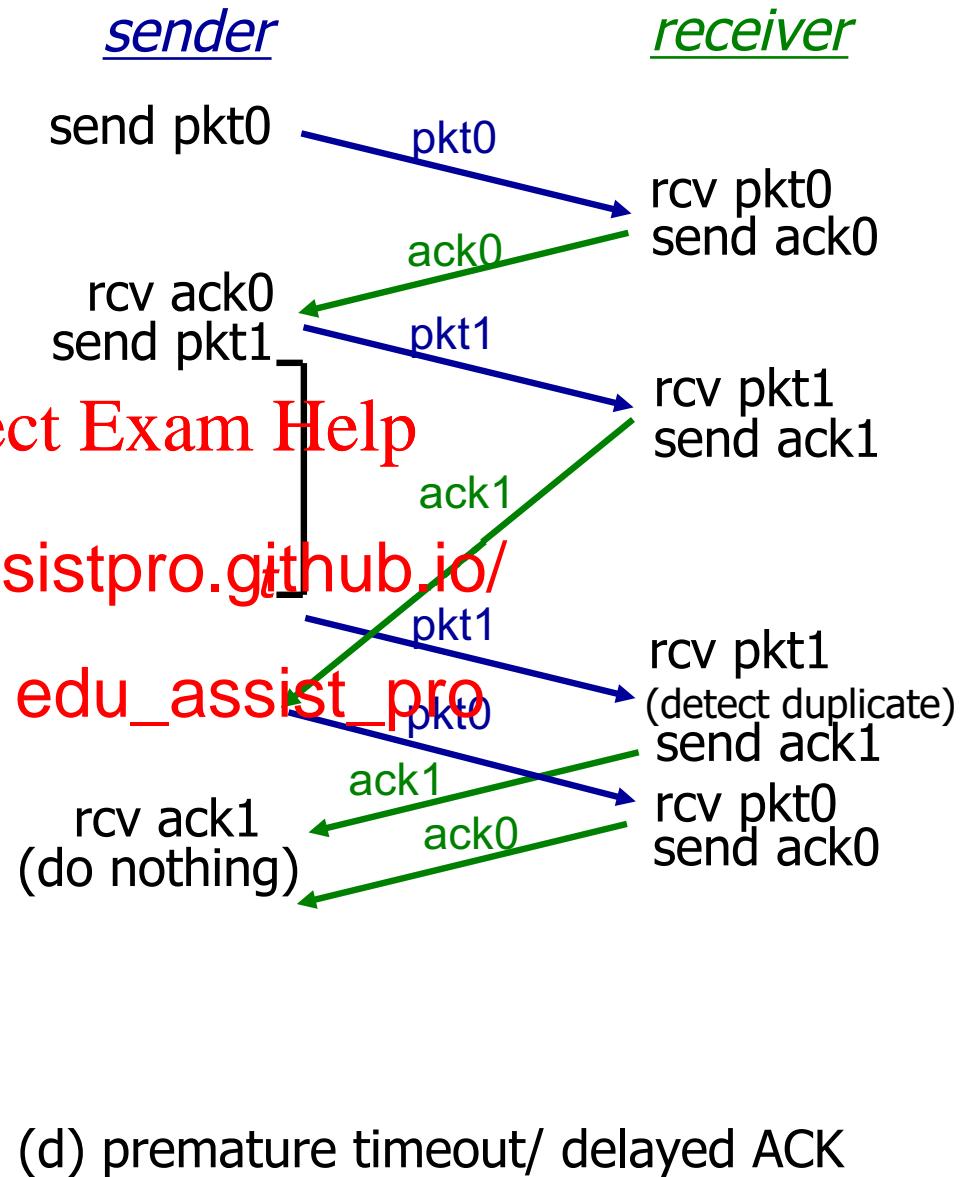
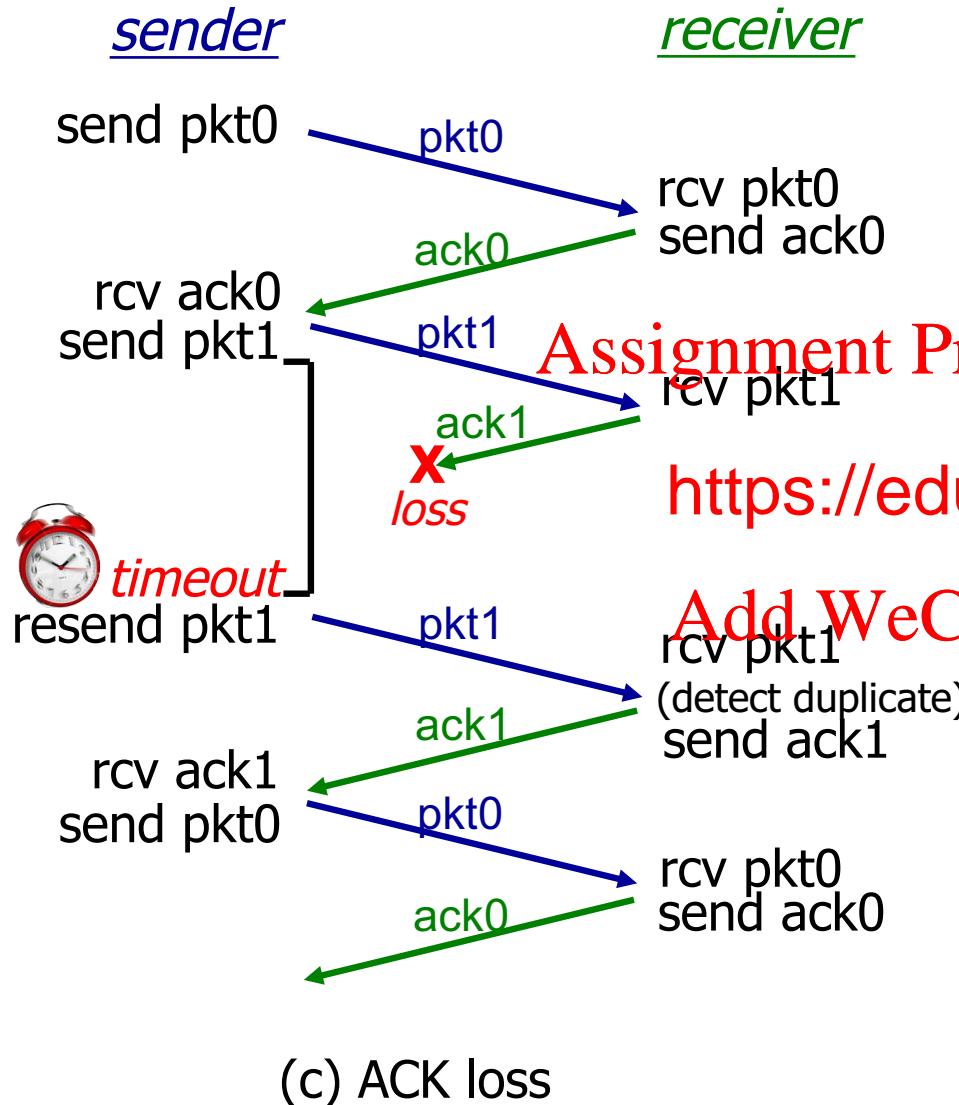
duplicate, but seq. #'s already handles this

- receiver must specify seq # of pkt being ACKed
- ❖ requires countdown timer

# rdt3.0 in action



# rdt3.0 in action

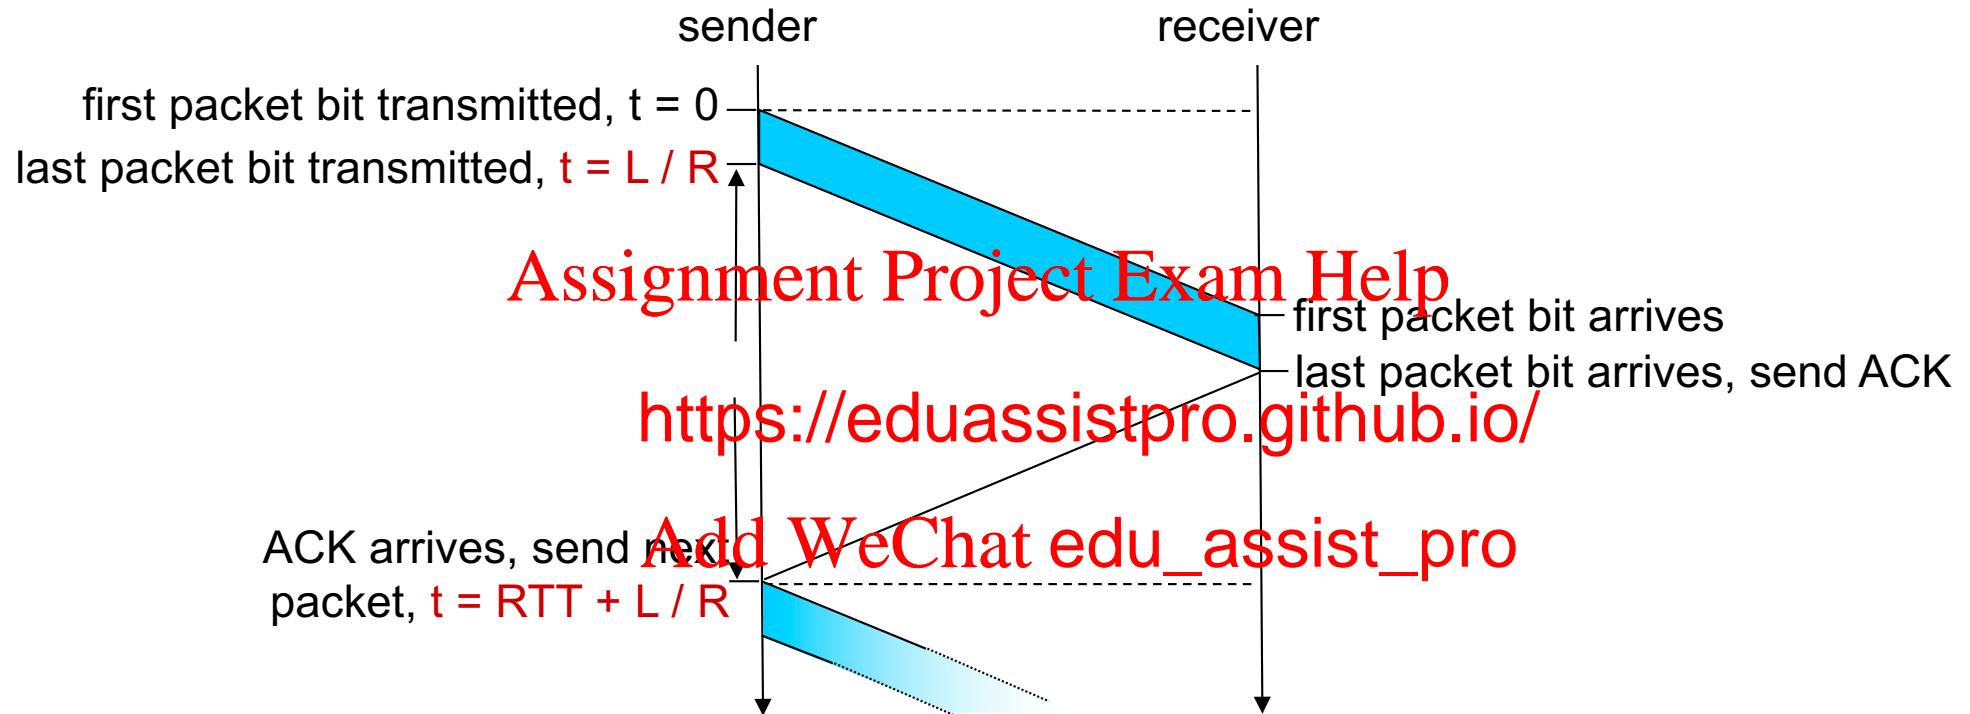


Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# rdt3.0: stop-and-wait operation



$$U_{\text{sender}} = \frac{L/R}{RTT + L/R}$$

# Performance of rdt3.0

- rdt3.0 is correct, but performance stinks
- e.g.: 1 Gbps link, 8000 bit packet and 30msec RTT:

Assignment Project Exam Help  
 $D_{trans} = \frac{L}{R}$      $\frac{8000 \text{ bits}}{8 \text{ microsecs}}$

<https://eduassistpro.github.io/>

- $U_{\text{sender}}$ : utilization – fraction der busy sending

Add WeChat edu\_assist\_pro

$$U_{\text{sender}} = \frac{L/R}{RTT + L/R} = \frac{0.008}{30.008} = 0.00027$$

- RTT=30 msec, 1KB pkt every 30.008 msec: 33kB/sec thruput over 1 Gbps link

- Network protocol limits use of physical resources!

# Pipelined protocols

**pipelining:** sender allows multiple, “in-flight”, yet-to-be-acknowledged pkts

- range of sequence numbers must be increased
- buffering at sender and/or receiver

Assignment Project Exam Help

<https://eduassistpro.github.io/>

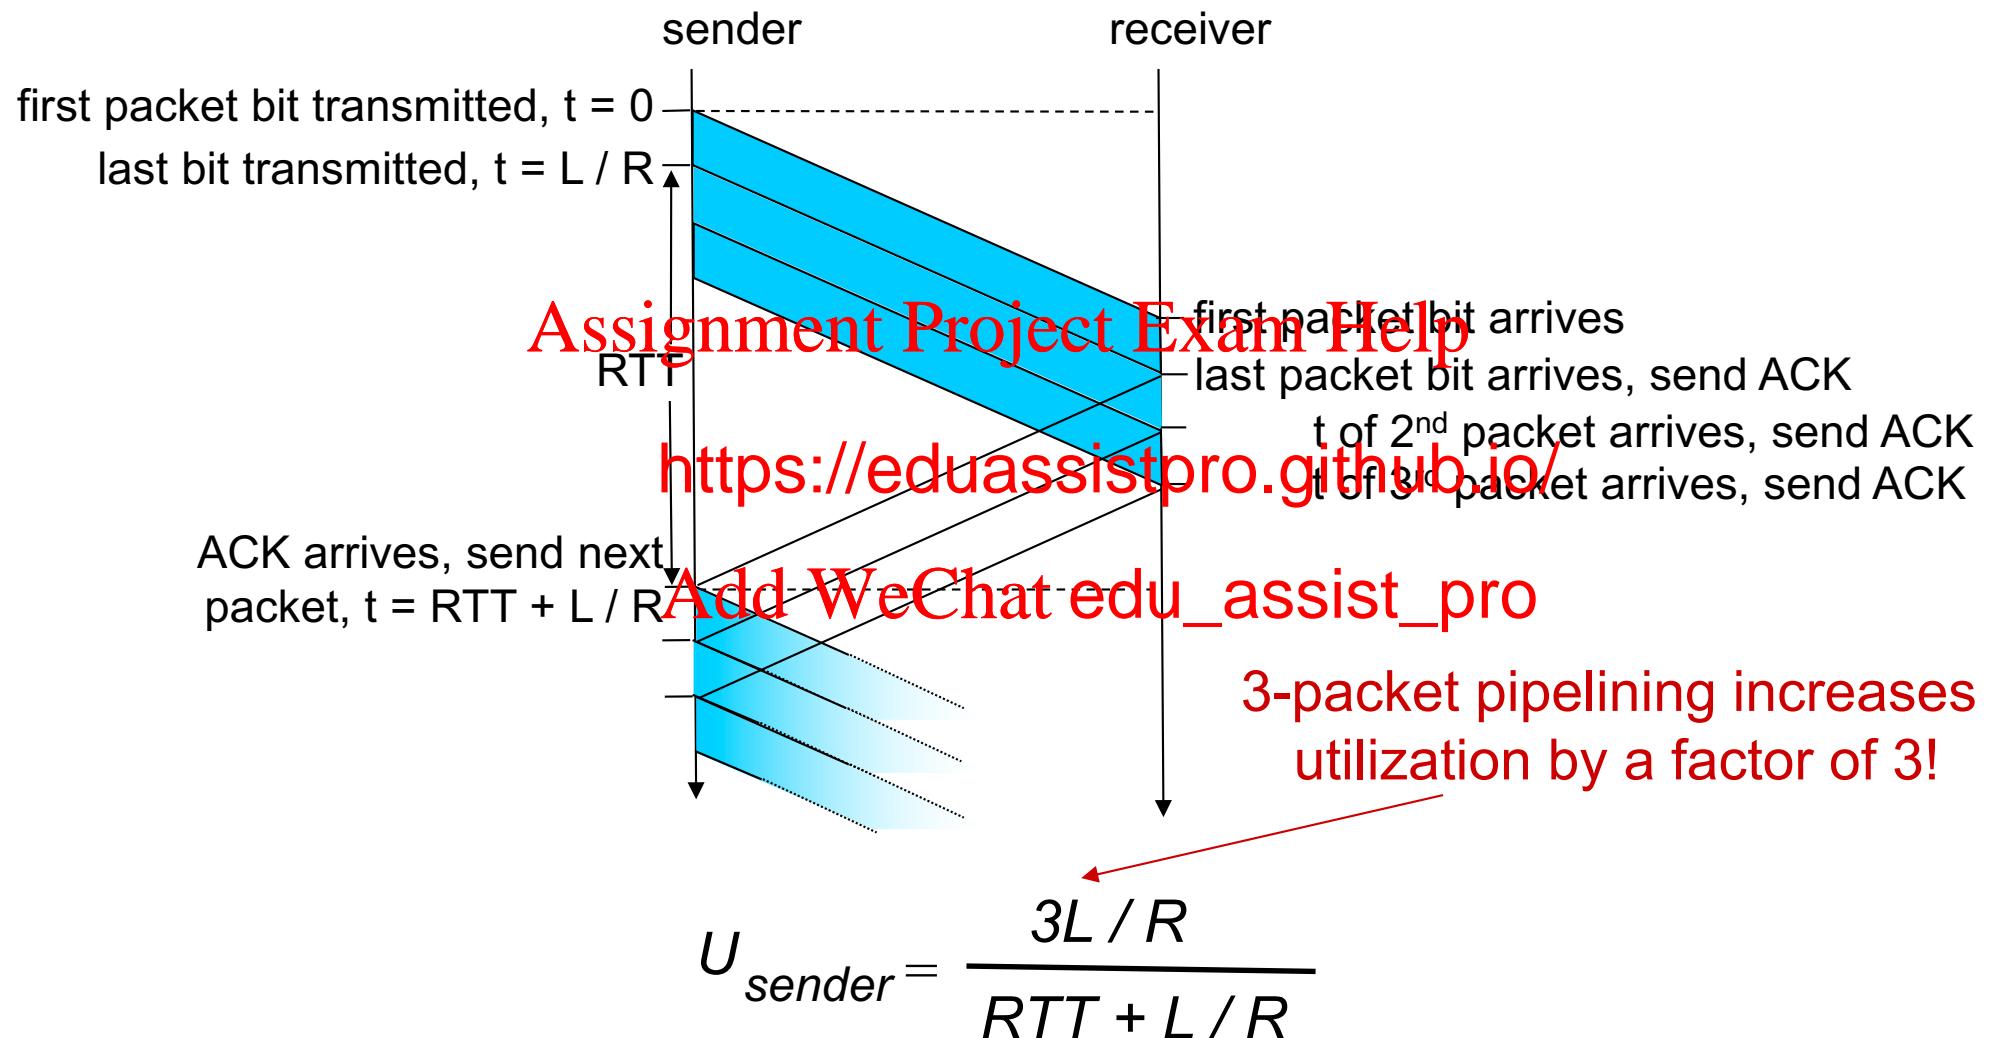


Add WeChat [edu\\_assist\\_pro](#)



- ❖ two generic forms of pipelined (sliding window) protocols: *go-Back-N, selective repeat*

# Pipelining: increased utilization



# Pipelined protocols: overview

## Go-Back-N:

- Sender can have up to N unacked packets in pipeline
- Sender has **single timer** for oldest unacked packet. When timer expires, retransmit unacked packets
- There is no buffer available at Receiver; out of order packets are discarded
- Receiver only sends **cumulative ack**, doesn't ack new packet if there's a gap

## Selective Repeat:

- Sender can have up to N unacked packets in pipeline
- Sender maintains timer for each packet, when timer expires, retransmit unacked packet
- Receiver has buffer, can accept out of order packets
- Receiver sends **individual ack** for each packet

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Go-Back-N: sender

- ❖ k-bit seq # in pkt header
- ❖ “window” of up to N, consecutive unack’ ed pkts allowed

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

- ❖ ACK(n):ACKs all pkts up to, i  $q \# n$  - “*cumulative ACK*”
  - may receive duplicate ACKs (see receiver)
- ❖ timer for oldest in-flight pkt
- ❖  $timeout(n)$ : retransmit packet n and all higher seq # pkts in window

Applets: [http://media.pearsoncmg.com/aw/aw\\_kurose\\_network\\_2/applets/go-back-n/go-back-n.html](http://media.pearsoncmg.com/aw/aw_kurose_network_2/applets/go-back-n/go-back-n.html)  
[http://www.ccs-labs.org/teaching/rn/animations/gbn\\_sr/](http://www.ccs-labs.org/teaching/rn/animations/gbn_sr/)

# GBN: sender extended FSM

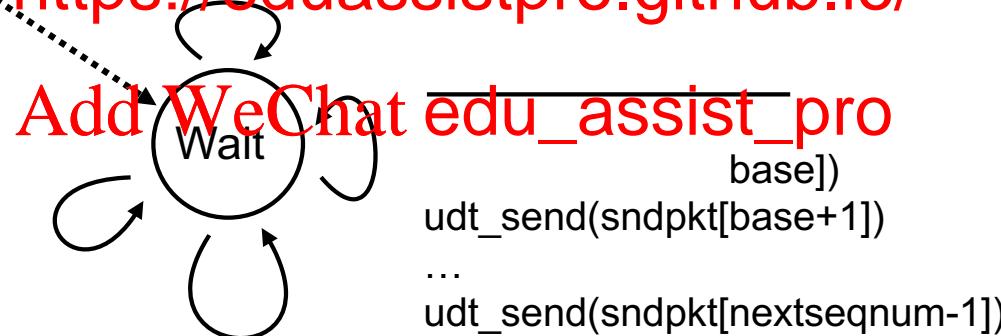
rdt\_send(data)

```
if (nextseqnum < base+N) {  
    sndpkt[nextseqnum] = make_pkt(nextseqnum,data,chksum)  
    udt_send(sndpkt[nextseqnum])  
    if (base == nextseqnum)  
        start_timer  
    nextseqnum++
```

Assignment Project Exam Help

$\frac{\Lambda}{\text{base}=1}$   
 $\text{nextseqnum}=1$

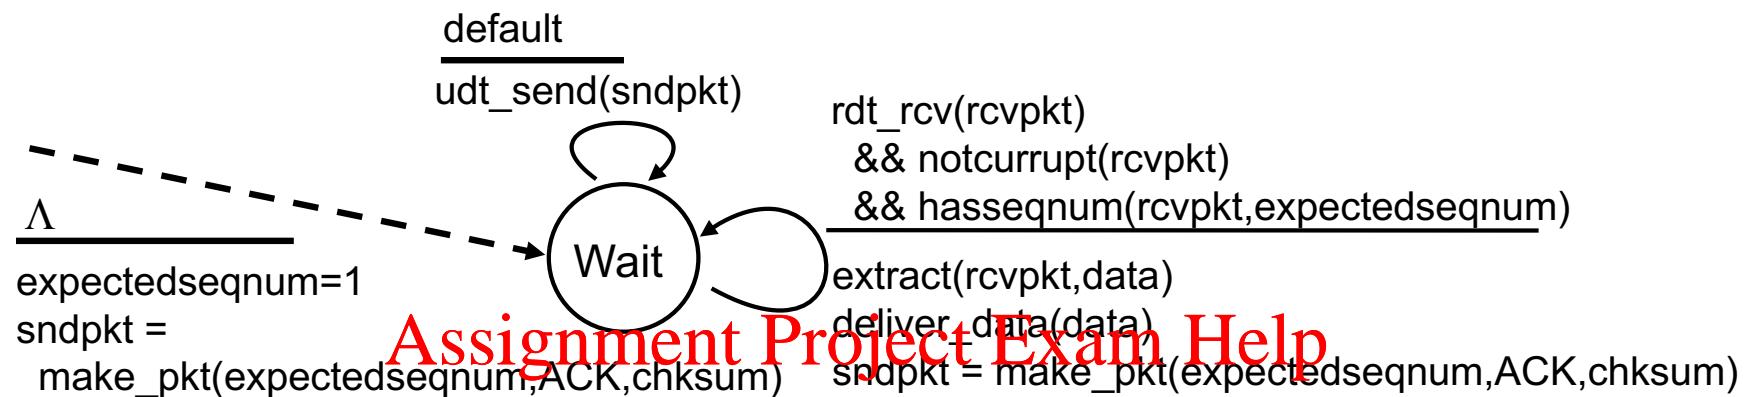
rdt\_rcv(rcvpkt)  
&& corrupt(rcvpkt)



rdt\_rcv(rcvpkt) &&  
notcorrupt(rcvpkt)

```
base = getacknum(rcvpkt)+1  
If (base == nextseqnum)  
    stop_timer  
else  
    start_timer
```

# GBN: receiver extended FSM



<https://eduassistpro.github.io/>

ACK-only: always send ACK for recently-received  
pkt with highest *in-order*

- may generate duplicate ACKs
- need only remember **expectedseqnum**
- ❖ out-of-order pkt:
  - discard (don't buffer): *no receiver buffering!*
  - re-ACK pkt with highest in-order seq #

# GBN in action

*sender window (N=4)*

0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8

0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8

0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8

*sender*

send pkt0  
send pkt1  
send pkt2  
send pkt3

(wait)

rcv ack0, s  
rcv ack1, s

send pkt2  
send pkt3  
send pkt4  
send pkt5

*receiver*

receive pkt0, send ack0  
receive pkt1, send ack1

receive pkt3, discard,  
(re)send ack1

receive pkt4, discard,  
(re)send ack1  
receive pkt5, discard,  
(re)send ack1

rcv pkt2, deliver, send ack2  
rcv pkt3, deliver, send ack3  
rcv pkt4, deliver, send ack4  
rcv pkt5, deliver, send ack5

<https://eduassistpro.github.io/>

ignore duplicate ACK

*pkt 2 timeout*



# Selective repeat

- ❖ receiver *individually* acknowledges all correctly received pkts
  - buffers pkts, as needed, for eventual in-order delivery to upper ~~Assignment Project Exam Help~~
- ❖ sender only receives ACKs for correctly received pkts
  - sender timer for each unACKed packet
- ❖ sender window
  - $N$  consecutive seq #'s
  - limits seq #'s of sent, unACKed pkts

Applet: [http://media.pearsoncmg.com/aw/aw\\_kurose\\_network\\_3/applets/SelectRepeat/SR.html](http://media.pearsoncmg.com/aw/aw_kurose_network_3/applets/SelectRepeat/SR.html)

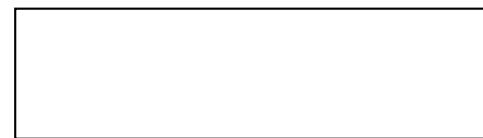
# Selective repeat: sender, receiver windows



Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro



# Selective repeat

## sender —

### **data from above:**

- ❖ if next available seq # in window, send pkt

### **timeout(n):**

- ❖ resend pkt n, re timer

### **ACK(n) in [sendbase,sendbase+N]:**

- ❖ mark pkt n as received
- ❖ if n smallest unACKed pkt, advance window base to next unACKed seq #

## receiver —

### **pkt n in [rcvbase, rcvbase+N-1]**

- ❖ send ACK(n)

❖ out-of-order: buffer  
order: deliver (also  
buffered, in-order  
window to

yet-received pkt

### **pkt n in [rcvbase-N,rcvbase-1]**

- ❖ ACK(n)

### **otherwise:**

- ❖ ignore

# Selective repeat in action

sender window ( $N=4$ )

0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8

0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8

0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8

sender

send pkt0  
send pkt1  
send pkt2  
send pkt3

(Wait)

Add WeChat edu\_assist\_pro

record ack3 arrived



*pkt 2 timeout*

send pkt2

record ack4 arrived

record ack5 arrived

receiver

receive pkt0, send ack0  
receive pkt1, send ack1

receive pkt3, buffer,  
send ack3

receive pkt4, buffer,  
send ack4  
receive pkt5, buffer,  
send ack5

rcv pkt2; deliver pkt2,  
pkt3, pkt4, pkt5; send ack2

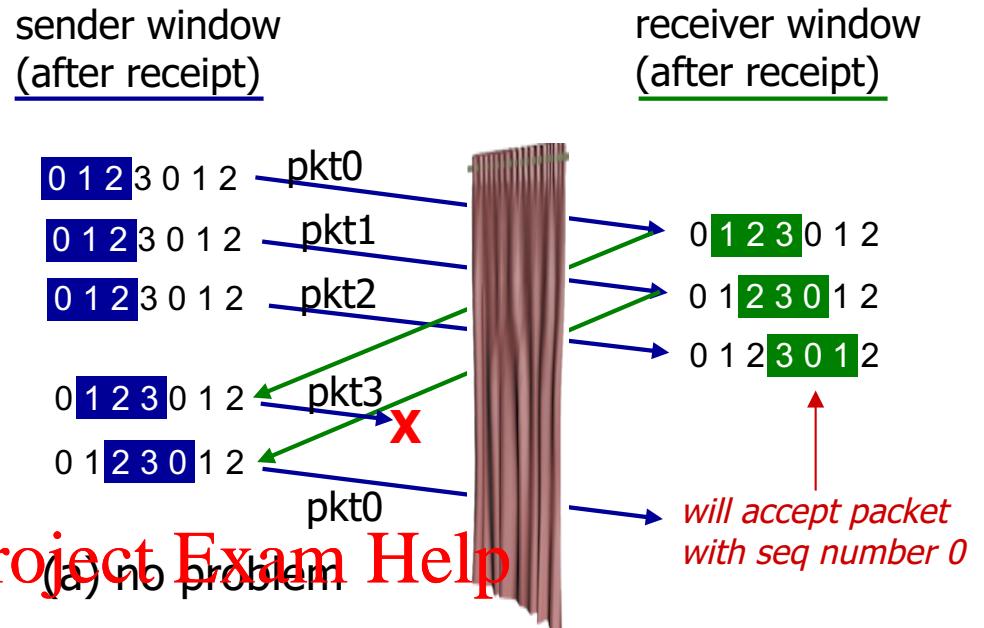
*Q: what happens when ack2 arrives?*

# Selective repeat: dilemma

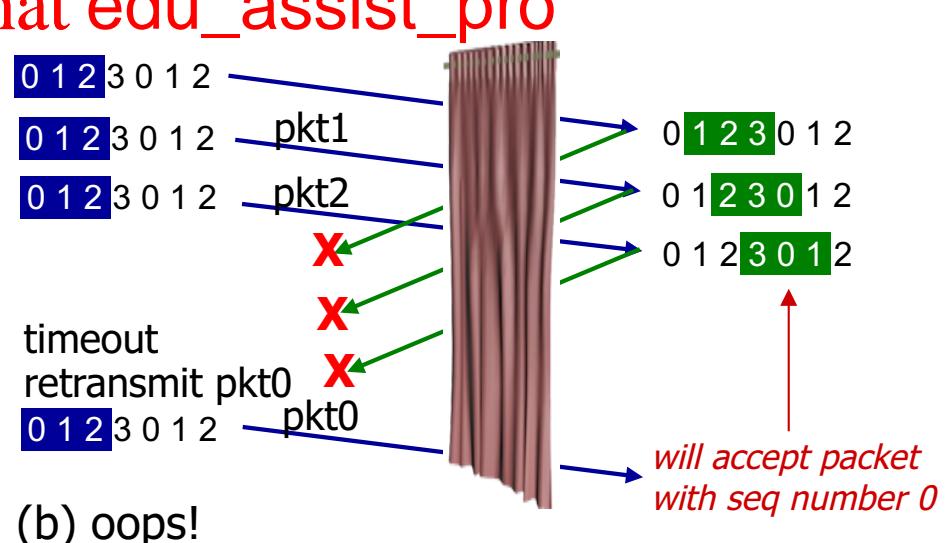
example:

- ❖ seq #'s: 0, 1, 2, 3
- ❖ window size=3
- ❖ receiver sees no difference in two scenarios!
- ❖ duplicate data accepted as new

Q: what relationship between seq # size and window size to avoid problem in (b)?



Assignment Project Exam Help  
https://eduassistpro.github.io/  
Viol identical in both cases!



# Recap: components of a solution

- ❖ Checksums (for error detection)
  - ❖ Timers (for loss detection)
  - ❖ Acknowledgments
    - cumulative
    - selective
  - ❖ Sequence numbers
  - ❖ Sliding Windows (for efficiency)
- Assignment Project Exam Help  
<https://eduassistpro.github.io/>  
Add WeChat edu\_assist\_pro
- ❖ Reliability protocols use the above to decide when and what to retransmit or acknowledge

# Transport Layer Outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

■ segment structure  
able data transfer

https://eduassistpro.github.io/  
Add WeChat edu\_assist\_pro  
3.6 mechanisms of congestion control

3.7 TCP congestion control

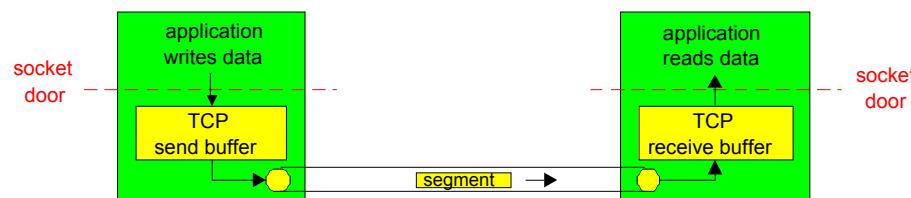
# Practical Reliability Questions

- ❖ How do the sender and receiver keep track of outstanding pipelined segments?
- ❖ How many segments should be pipelined?  
*Assignment Project Exam Help*
- ❖ How do we calculate sequence numbers?
- ❖ What does connection establishment and teardown look like?  
*Add WeChat edu\_assist\_pro  
<https://eduassistpro.github.io/>*
- ❖ How should we choose timeout values?

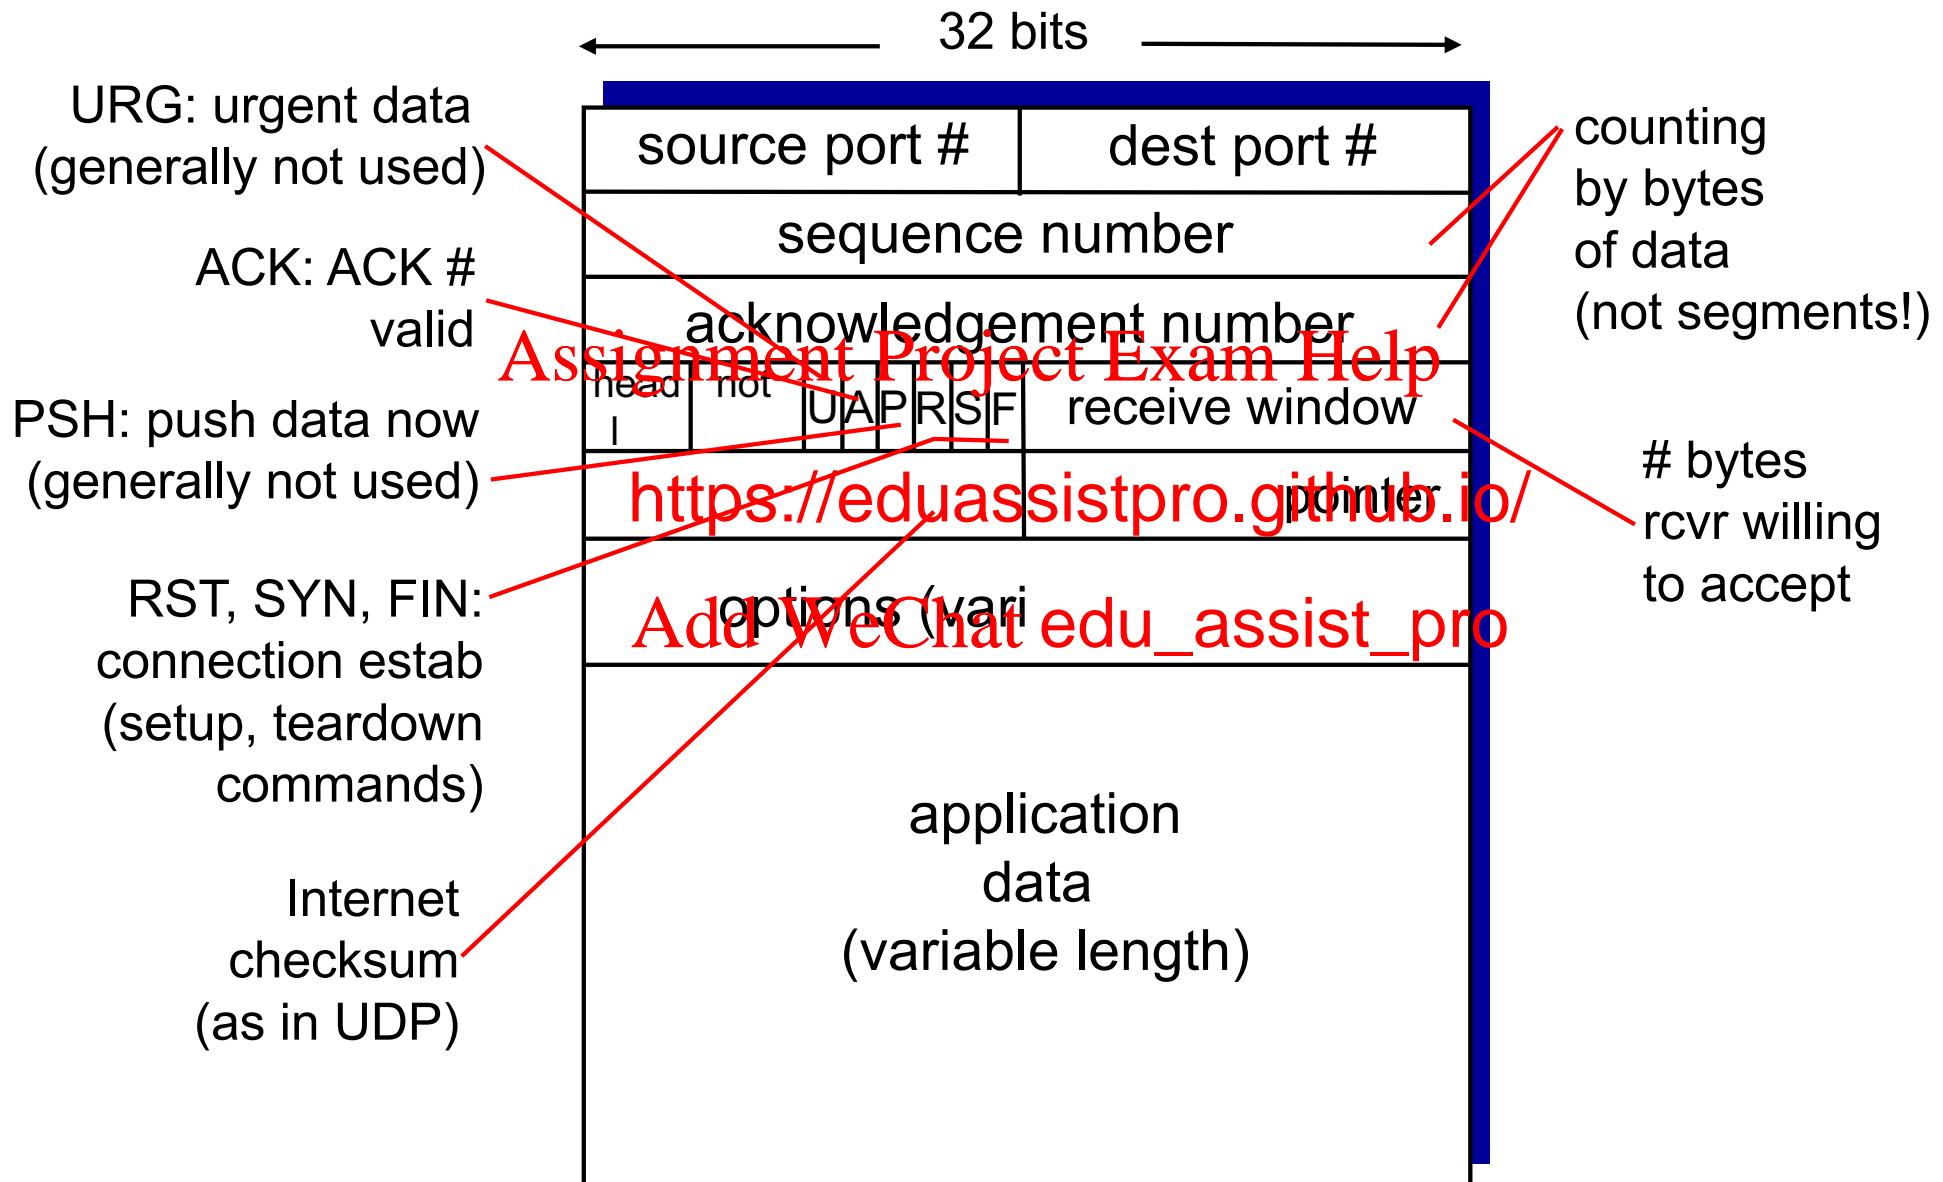
# TCP: Overview

RFCs: 793, 1122, 1323, 2018, 2581

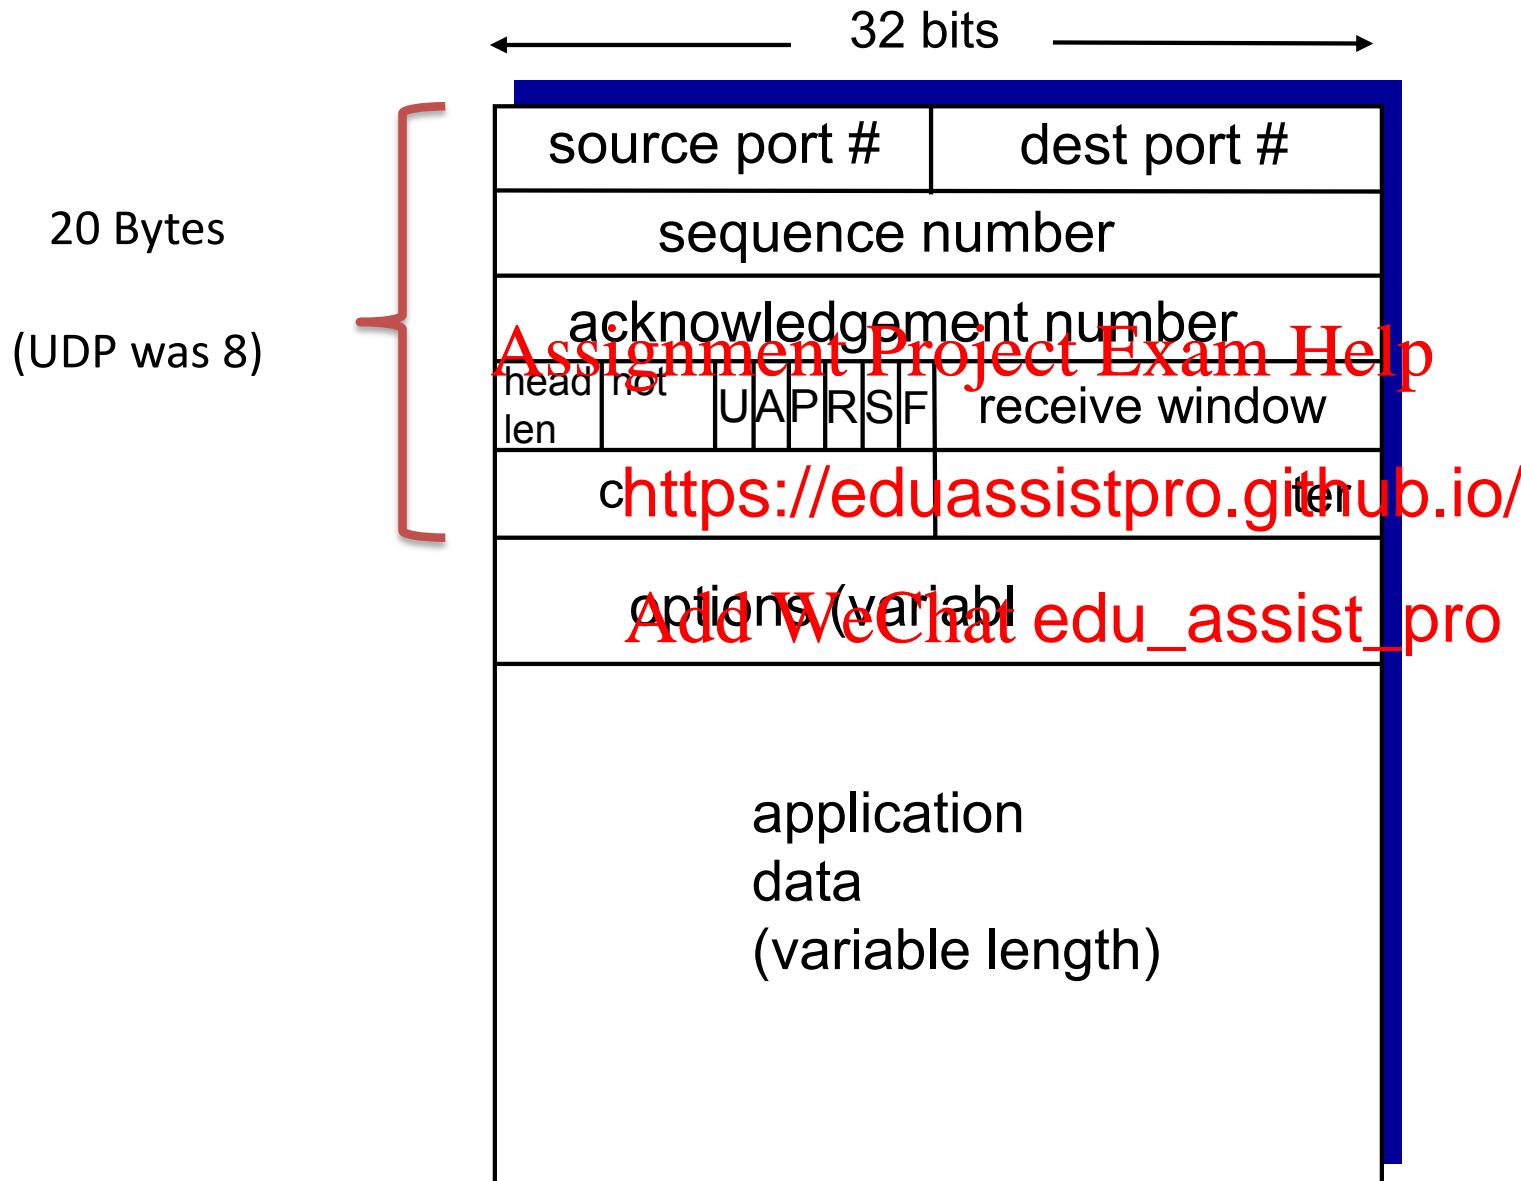
- ❖ **point-to-point:**
  - one sender, one receiver
- ❖ **reliable, in order byte stream:**
  - no “message
- ❖ **pipelined:**
  - TCP congestion and flow control set window size
- ❖ **send and receive buffers**
- ❖ **full duplex data:**
  - bi-directional data flow in same connection
  - MSS: maximum segment size
- ❖ **connection-oriented:**
  - three-way handshaking (exchange of control msgs) initiates sender, receiver state before data exchange
- ❖ **flow controlled:**
  - sender will not overwhelm receiver



# TCP segment structure



# TCP segment structure



# Transport Layer Outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

■ segment structure  
able data transfer

https://eduassistpro.github.io/  
Add WeChat edu\_assist\_pro  
ction management  
es of congestion  
control

3.7 TCP congestion control

# Recall: Components of a solution for reliable transport

---

- ❖ Checksums (for error detection)
- ❖ Timers (for loss detection)  
Assignment Project Exam Help
- ❖ Acknowledgments
  - cumulative
  - selective<https://eduassistpro.github.io/>  
Add WeChat edu\_assist\_pro
- ❖ Sequence numbers (duplicates, windows)
- ❖ Sliding Windows (for efficiency)
  - Go-Back-N (GBN)
  - Selective Repeat (SR)

# What does TCP do?

---

Many of our previous ideas, but some key differences

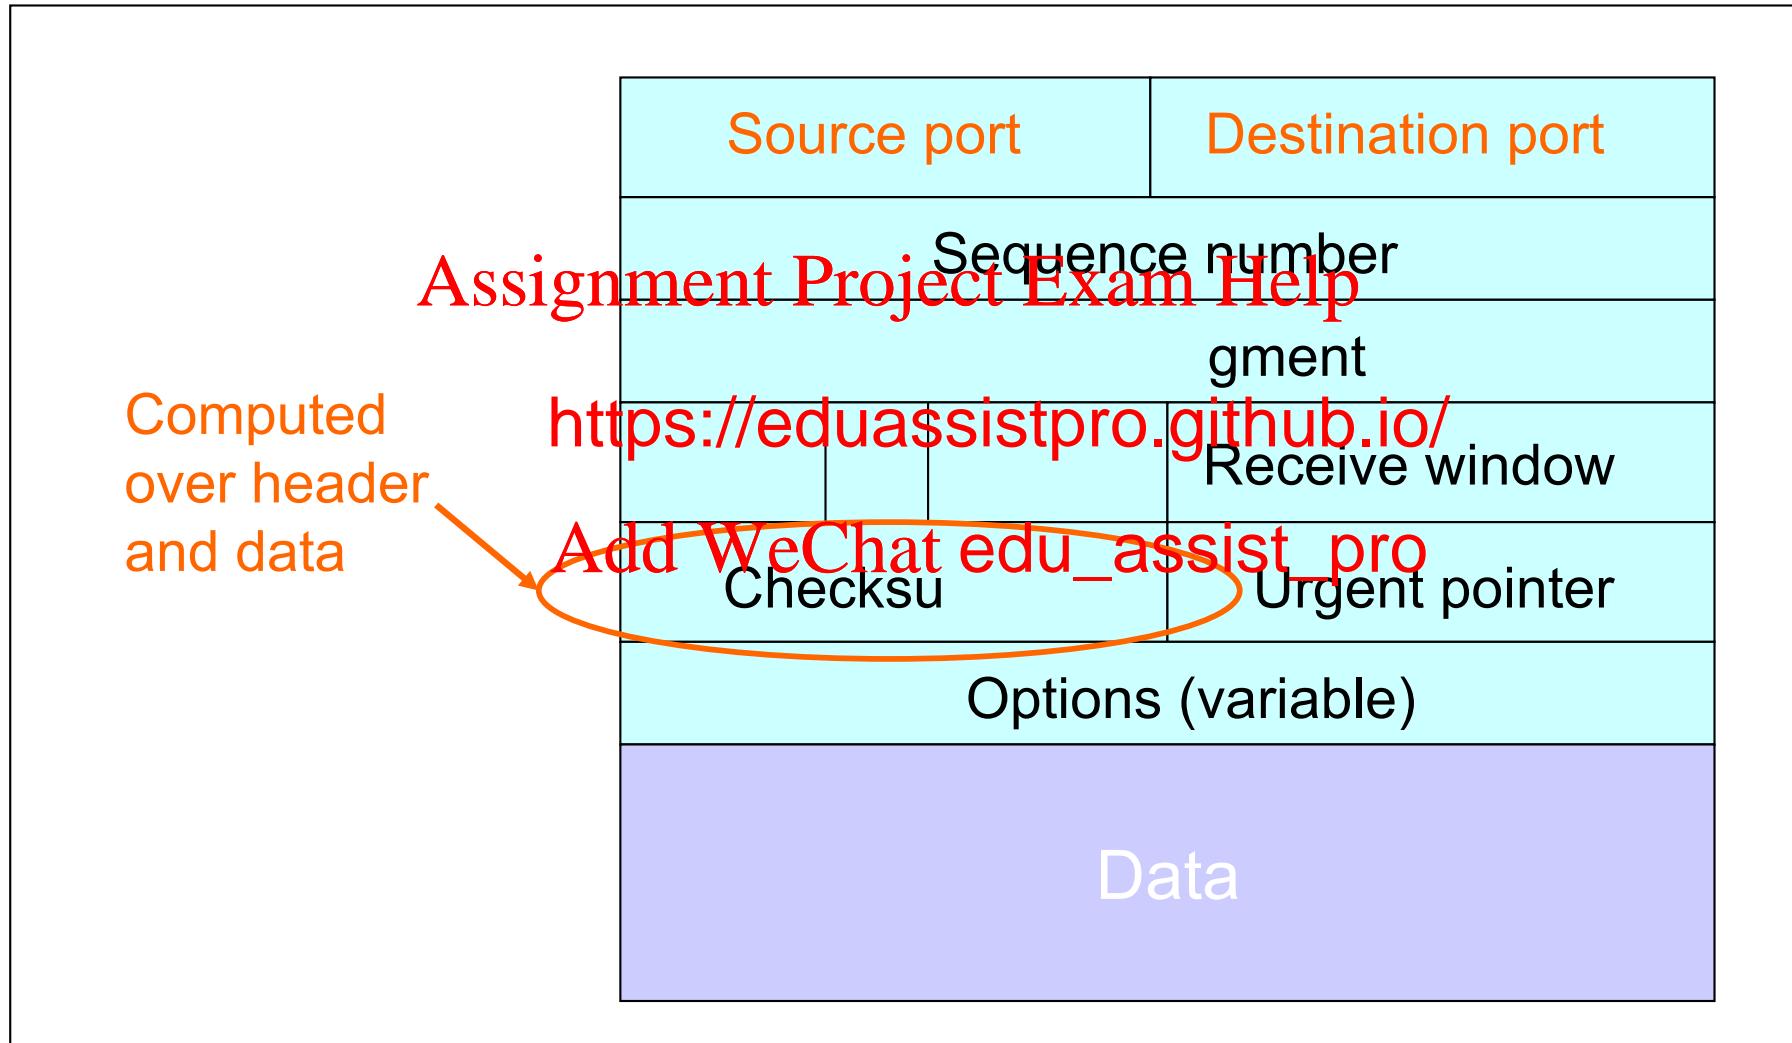
- ❖ Checksum

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# TCP Header



# What does TCP do?

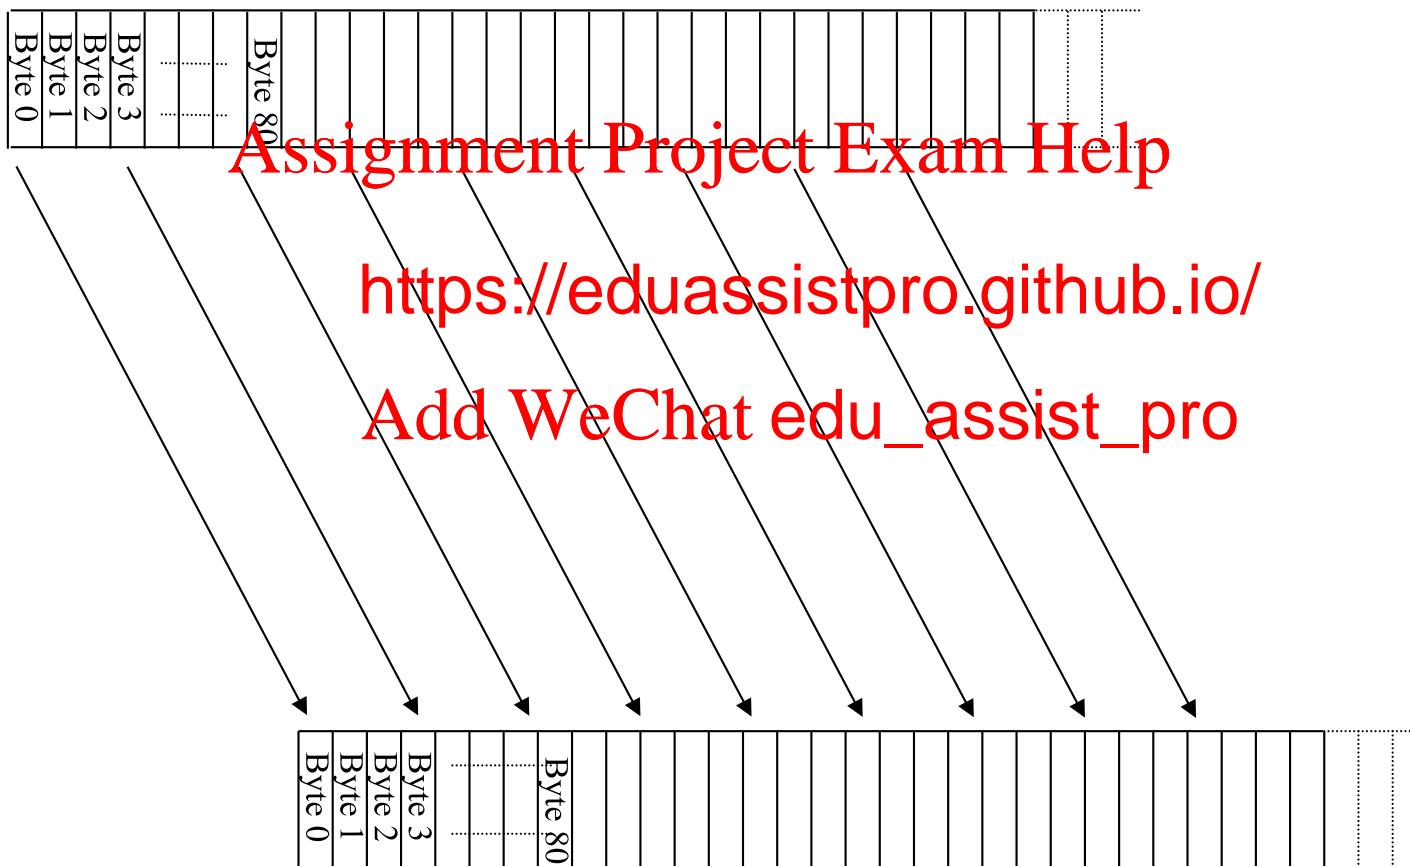
---

Many of our previous ideas, but some key differences

- ❖ Checksum [Assignment Project Exam Help](#)
- ❖ **Sequence numb** <https://eduassistpro.github.io/>  
[Add WeChat edu\\_assist\\_pro](#)

# TCP “Stream of Bytes” Service ..

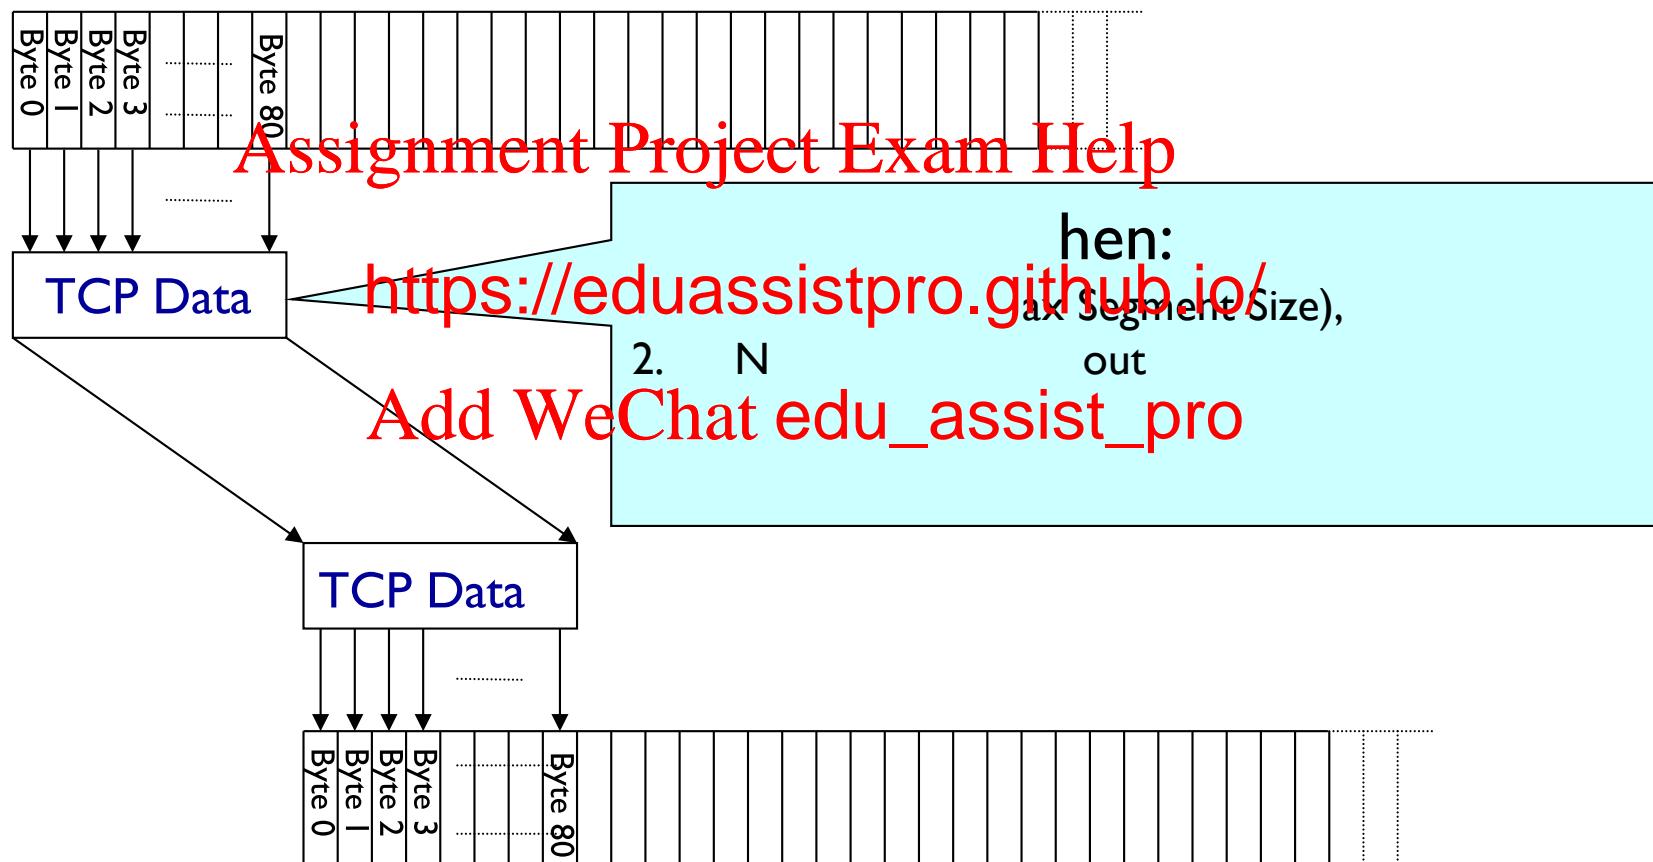
# Application @ Host A



# Application @ Host B

# .. Provided Using TCP “Segments”

Host A



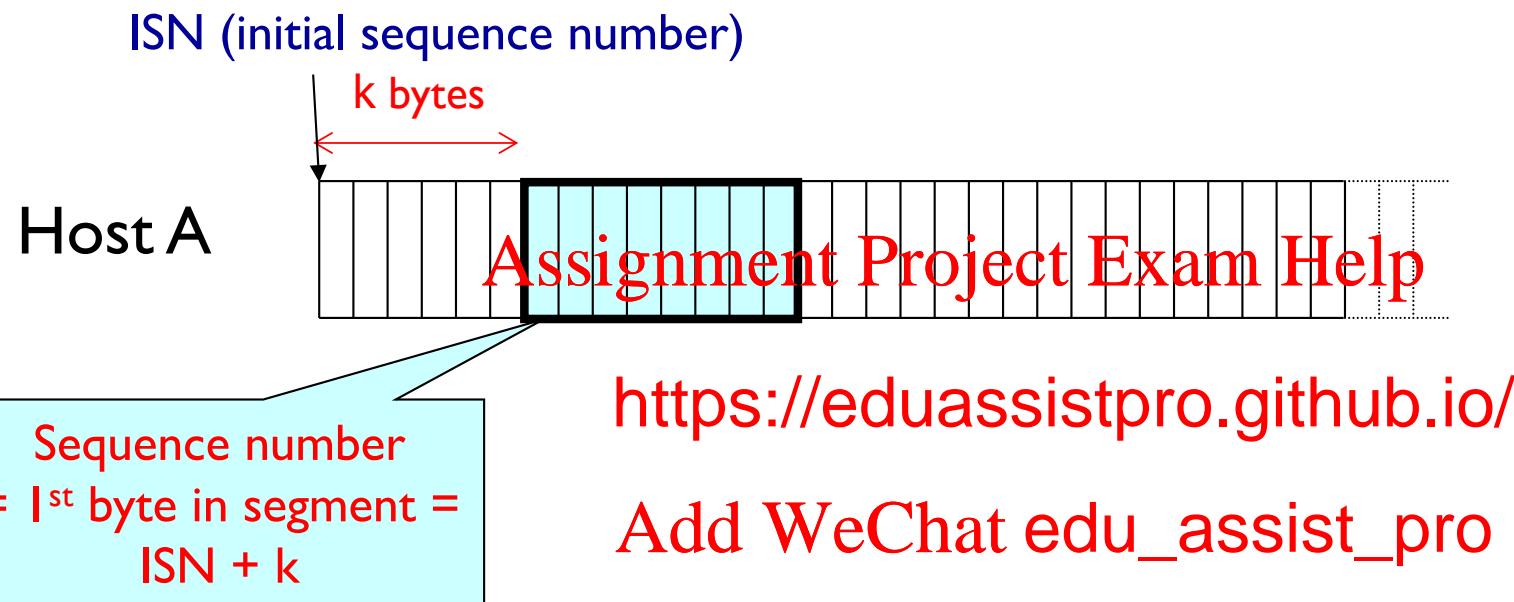
Host B

# TCP Segment Size



- ❖ IP packet Assignment Project Exam Help
  - No bigger than MTU Unit (MTU)
  - E.g., up to 150 <https://eduassistpro.github.io/>
- ❖ TCP packet Add WeChat edu\_assist\_pro
  - IP packet with a TCP header and data inside
  - TCP header  $\geq$  20 bytes long
- ❖ TCP segment
  - No more than Maximum Segment Size (MSS) bytes
  - E.g., up to 1460 consecutive bytes from the stream
  - $MSS = MTU - (\text{IP header}) - (\text{TCP header})$

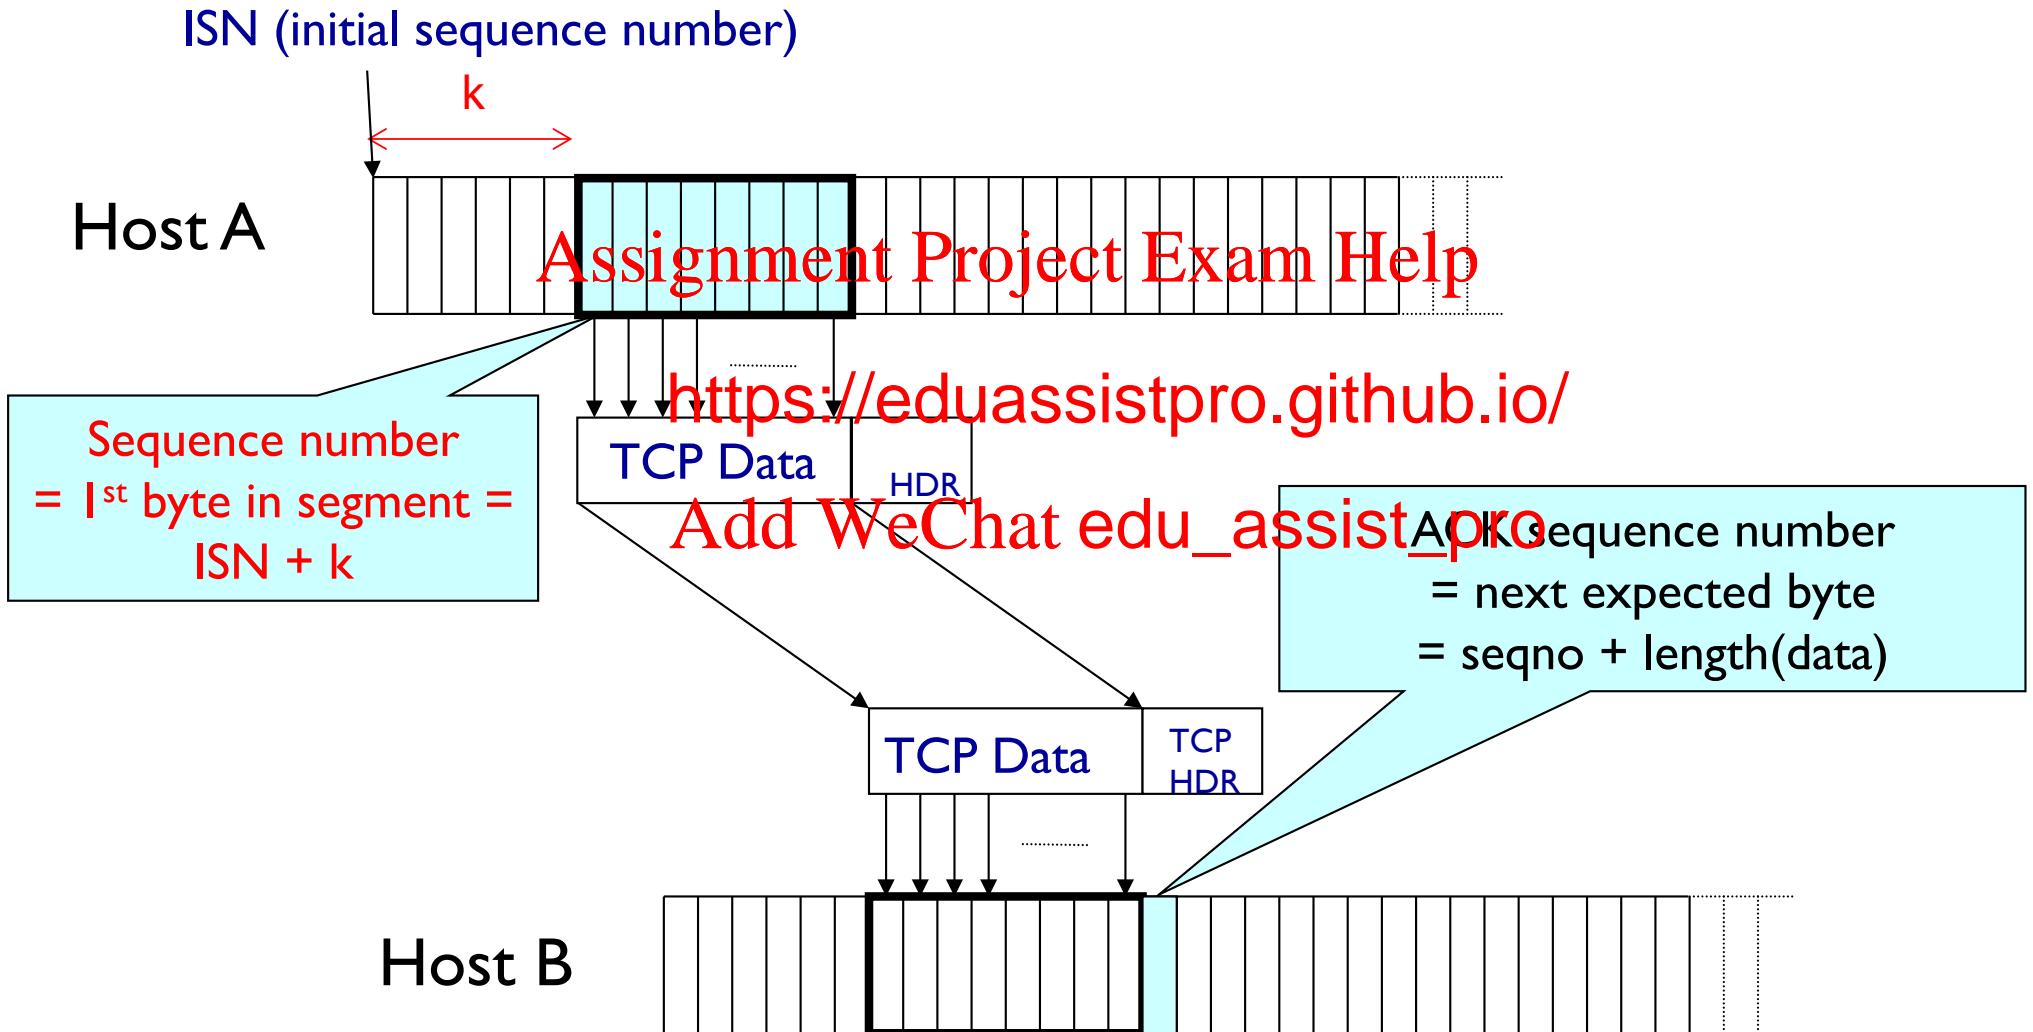
# Sequence Numbers



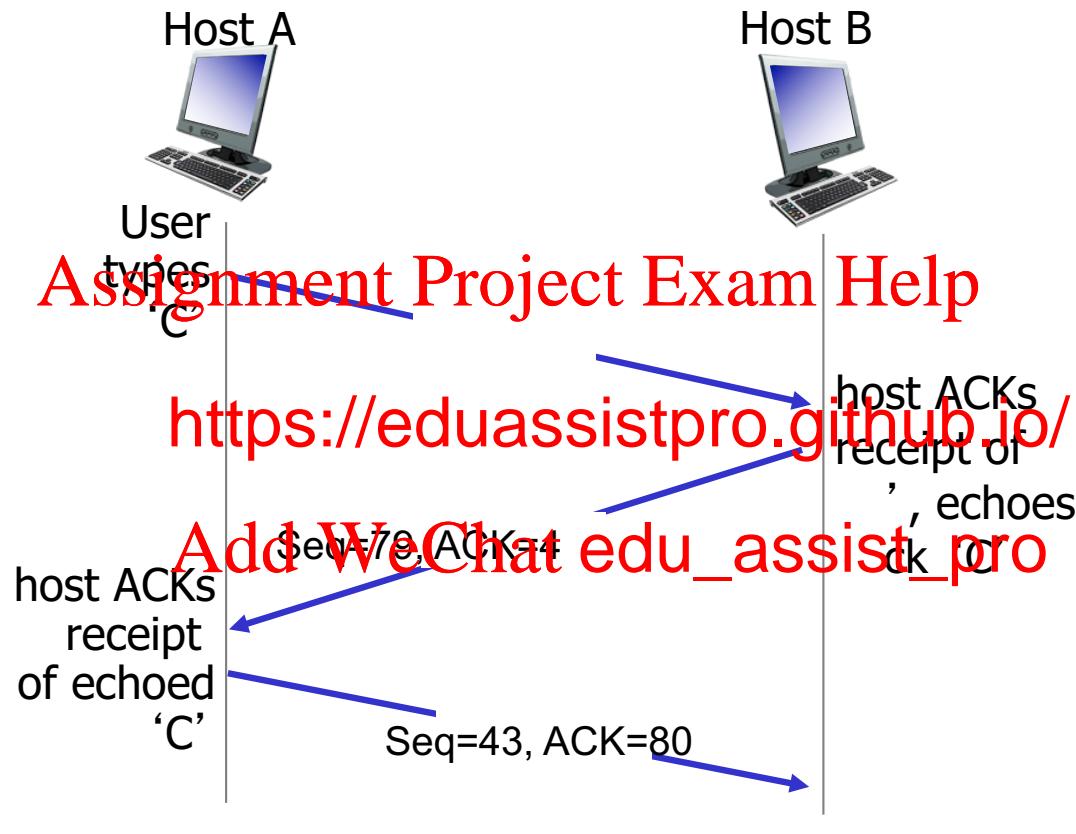
## Sequence numbers:

- byte stream “number” of first byte in segment’s data

# Sequence & Ack Numbers



# TCP seq. numbers, ACKs



simple telnet scenario

# What does TCP do?

Most of our previous tricks, but a few differences

- ❖ Checksum
- ❖ Sequence numbers are byte offsets
- ❖ Receiver sends cumulative acknowledgments (like GBN)

Assignment Project Exam Help

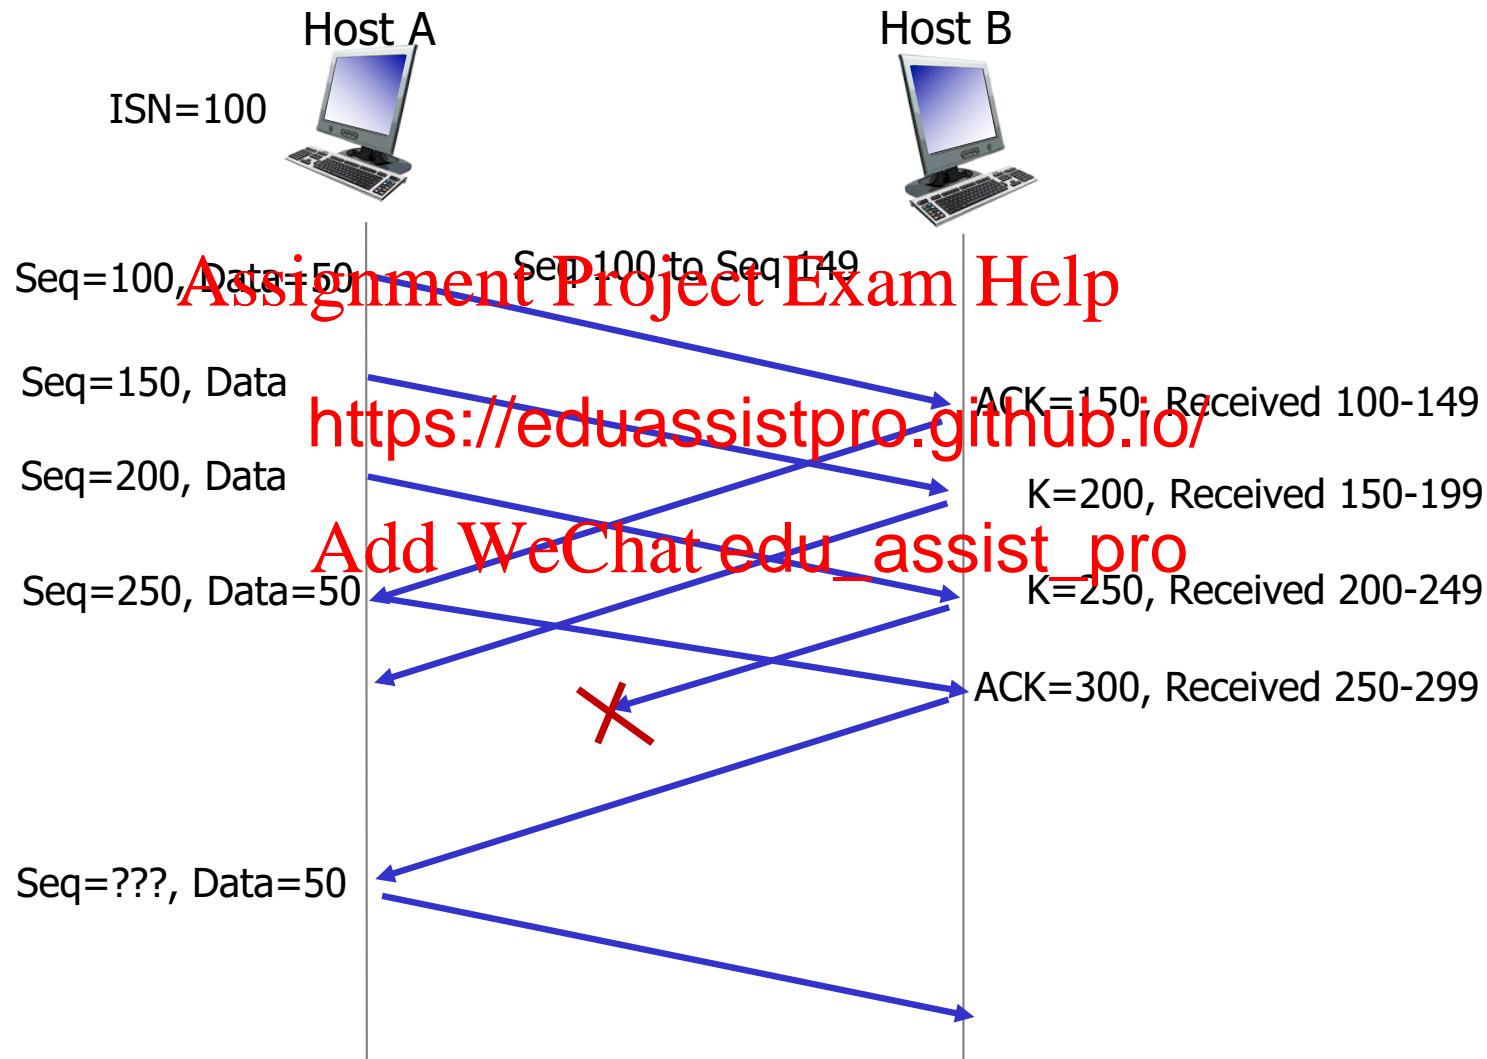
<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

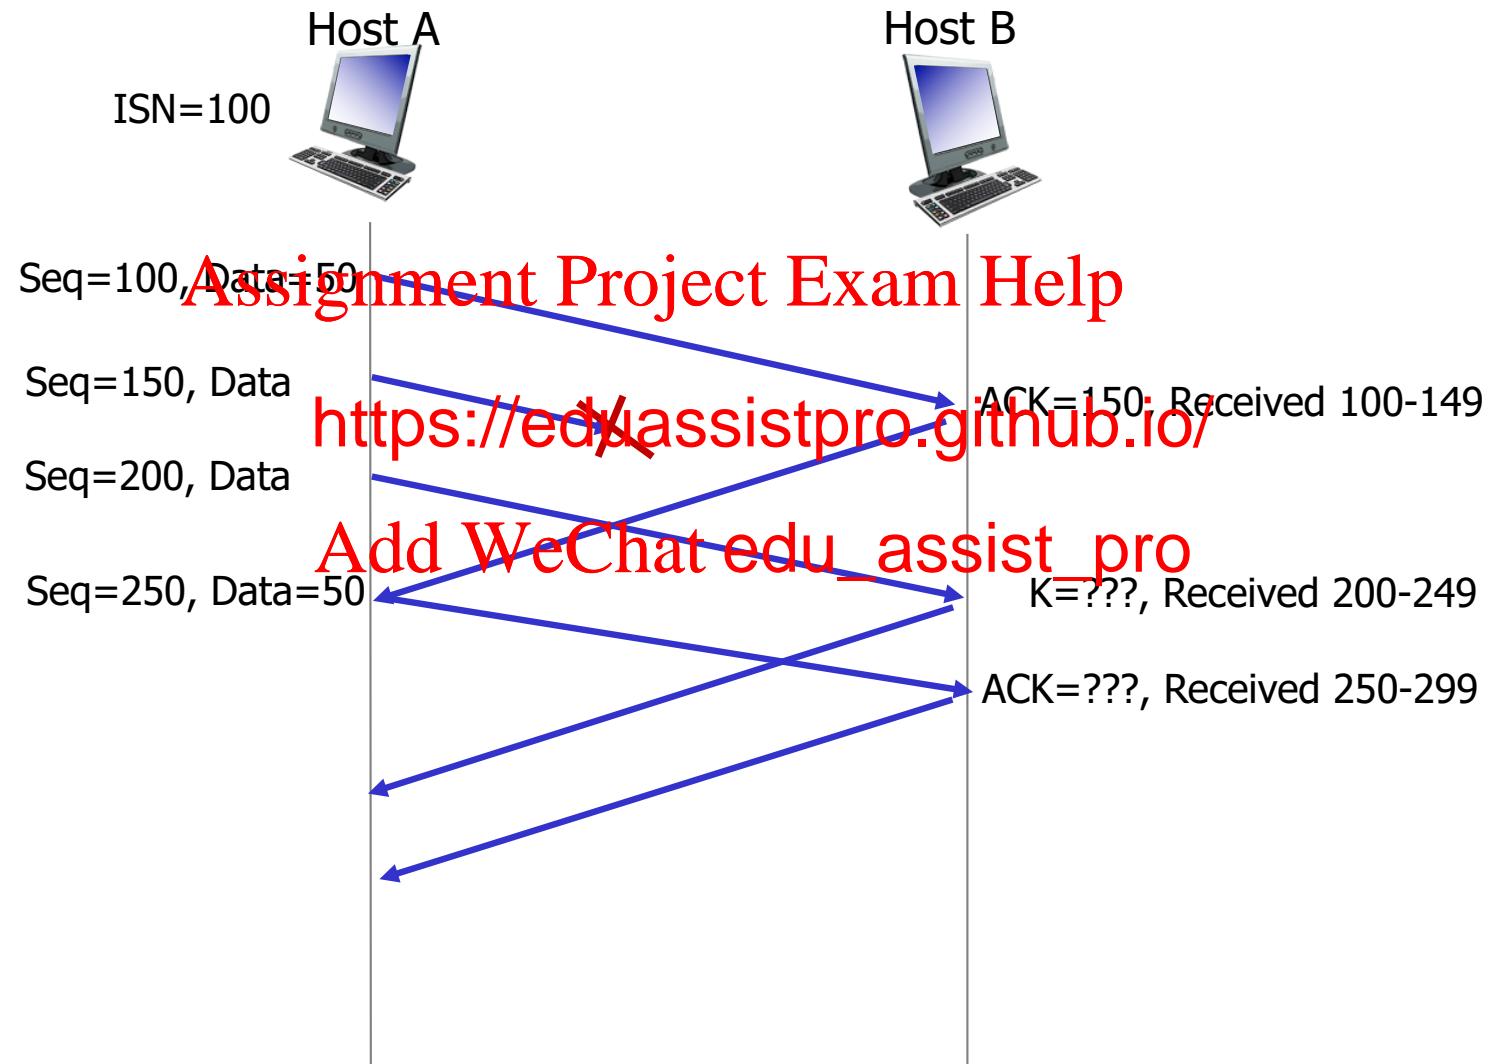
# ACKing and Sequence Numbers

- ❖ Sender sends packet
  - Data starts with sequence number X
  - Packet contains B bytes [X, X+1, X+2, ..., X+B-1]
- ❖ Upon receipt of a packet
  - If all data prior to the current byte received is in-order
    - ACK acknowledges  $X+B$  (because it expected byte  $X+B-1 < X$ )
  - If highest in-order byte received is  $Y$ 
    - ACK acknowledges  $Y+1$
    - Even if this has been ACKed before

# TCP seq. numbers, ACKs



# TCP seq. numbers, ACKs



# Normal Pattern

---

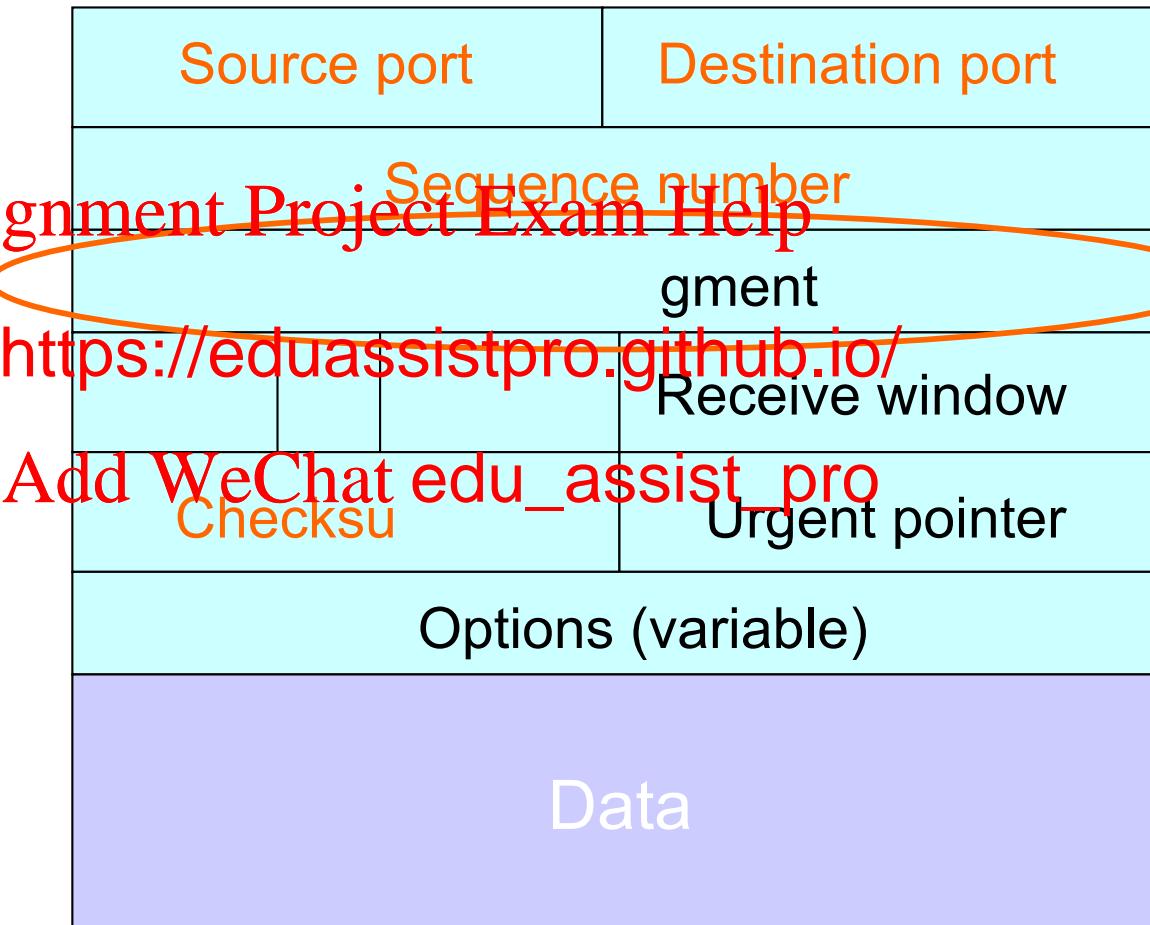
- ❖ Sender: seqno=X, length=B
- ❖ Receiver: ACK=X+B
- ❖ Sender: seqno=X+B, length=B
- ❖ Receiver: ACK=X+2B
- ❖ Sender: seqno= <https://eduassistpro.github.io/>
- ❖ Seqno of next packet is same as Add WeChat edu\_assist\_pro field

# Packet Loss

- ❖ Sender: seqno=X, length=B
- ❖ Receiver: ACK=X+B
- ❖ Sender: ~~seqno=X+B, length=B~~ LOST  
**Assignment Project Exam Help**
- ❖ Sender: seqno= <https://eduassistpro.github.io/>
- ❖ Receiver: ACK = X+B  
**Add WeChat edu\_assist\_pro**

# TCP Header

Acknowledgment gives seqno just beyond highest seqno received **in order** (“*What Byte is Next*”)



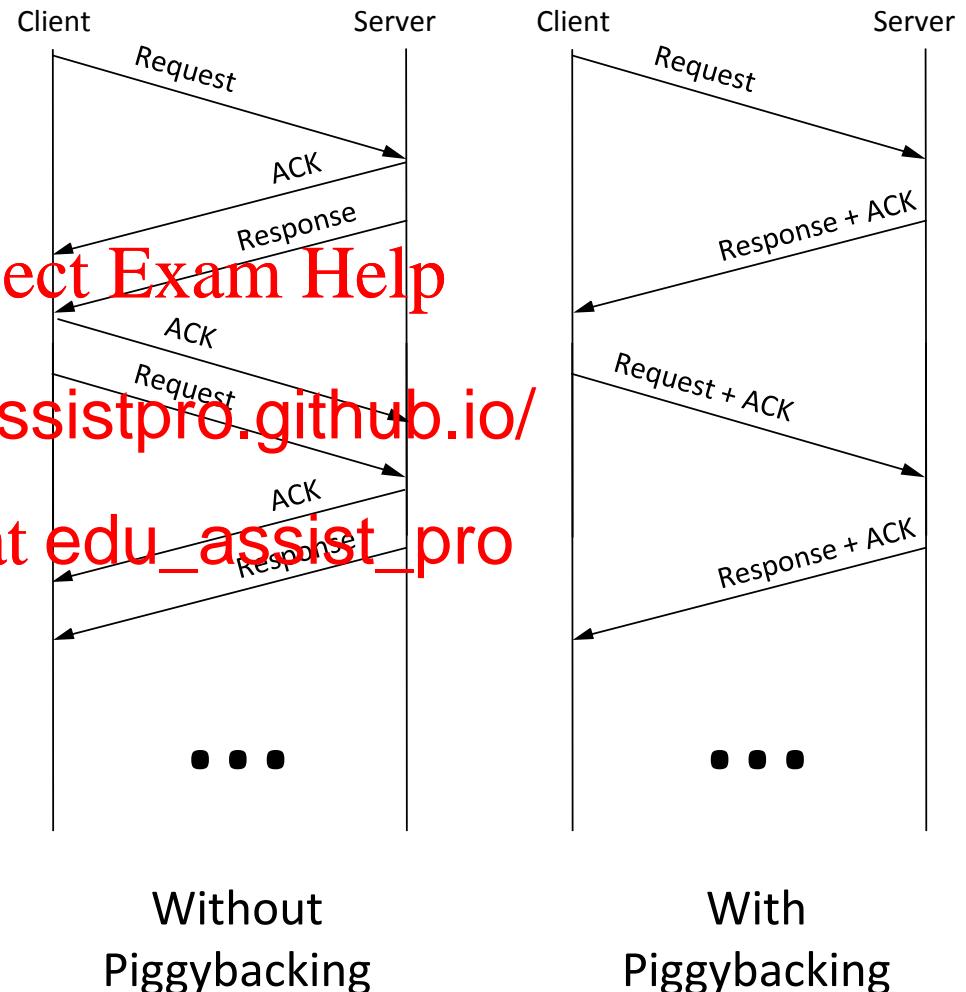
# Piggybacking

- ❖ So far, we've assumed distinct “sender” and “receiver” roles

Assignment Project Exam Help

- ❖ In reality, usu <https://eduassistpro.github.io/> sides of a connection send some data

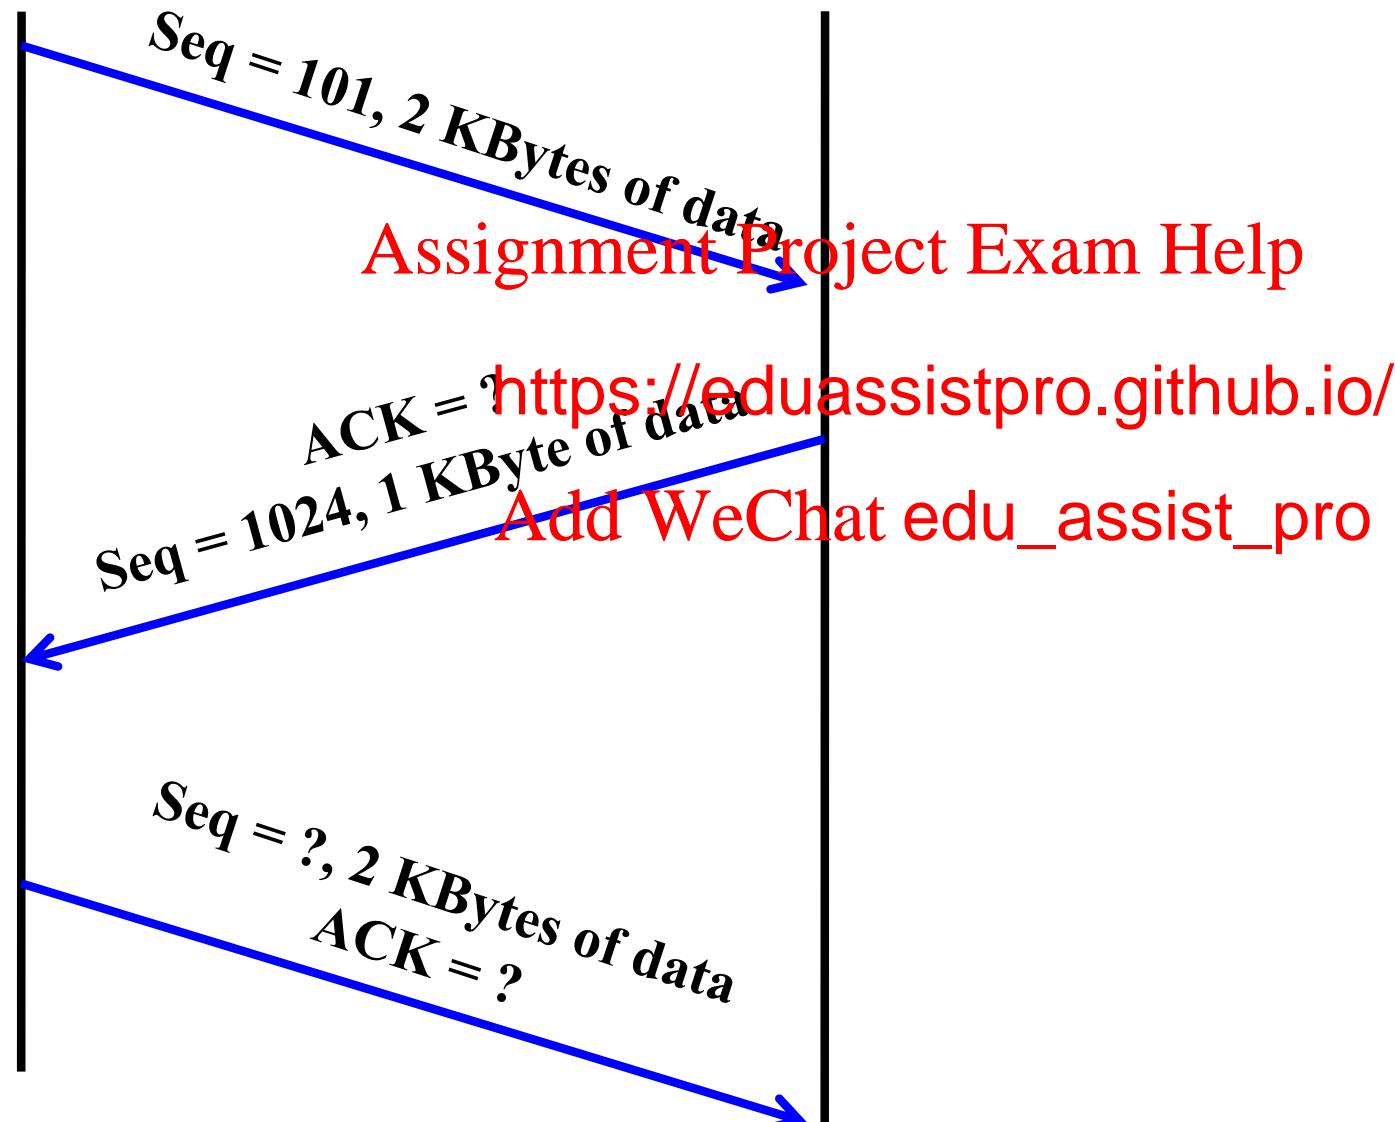
Add WeChat edu\_assist\_pro



Without  
Piggybacking

With  
Piggybacking

# Quiz



# What does TCP do?

Most of our previous tricks, but a few differences

- ❖ Checksum
- ❖ Sequence numbers are byte offsets
- ❖ Receiver sends cumulative acknowledgments (like GBN)
- ❖ Receivers can buffer segments (like SR)  
<https://eduassistpro.github.io/>  
Add WeChat edu\_assist\_pro

# Loss with cumulative ACKs

---

- ❖ Sender sends packets with 100Bytes and sequence numbers:
  - 100, 200, 300, 400, 500, 600, 700, 800, 900, ...  
<https://eduassistpro.github.io/>
- ❖ Assume the ~~first packet~~ ~~Add WeChat edu\_assist\_500~~ (500) is lost, but no others
- ❖ Stream of ACKs will be:
  - 200, 300, 400, 500, 500, 500, 500, ...

# What does TCP do?

Most of our previous tricks, but a few differences

- ❖ Checksum
- ❖ Sequence numbers are byte offsets
- ❖ Receiver sends cumulative acknowledgments (like GBN)
- ❖ Receivers do not duplicate ACKs (like SR)
- ❖ Sender maintains a single retransmit buffer (like GBN) and retransmits on timeout

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat `edu_assist_pro` (like GBN) and  
retransmits on timeout

# TCP round trip time, timeout

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# TCP round trip time, timeout

Q: how to set TCP timeout value?

- ❖ longer than RTT
  - but RTT vari
- ❖ too short: pre timeout, unnecessary retransmissions
- ❖ too long: slow reaction to segment loss and connection has lower throughput

Q: how to estimate RTT?

- ❖ **SampleRTT**: measured time from segment transmission until ACK

<https://eduassistpro.github.io/>

Add WeChat `edu_assist_pro` RTT will vary, want RTT “smoother”

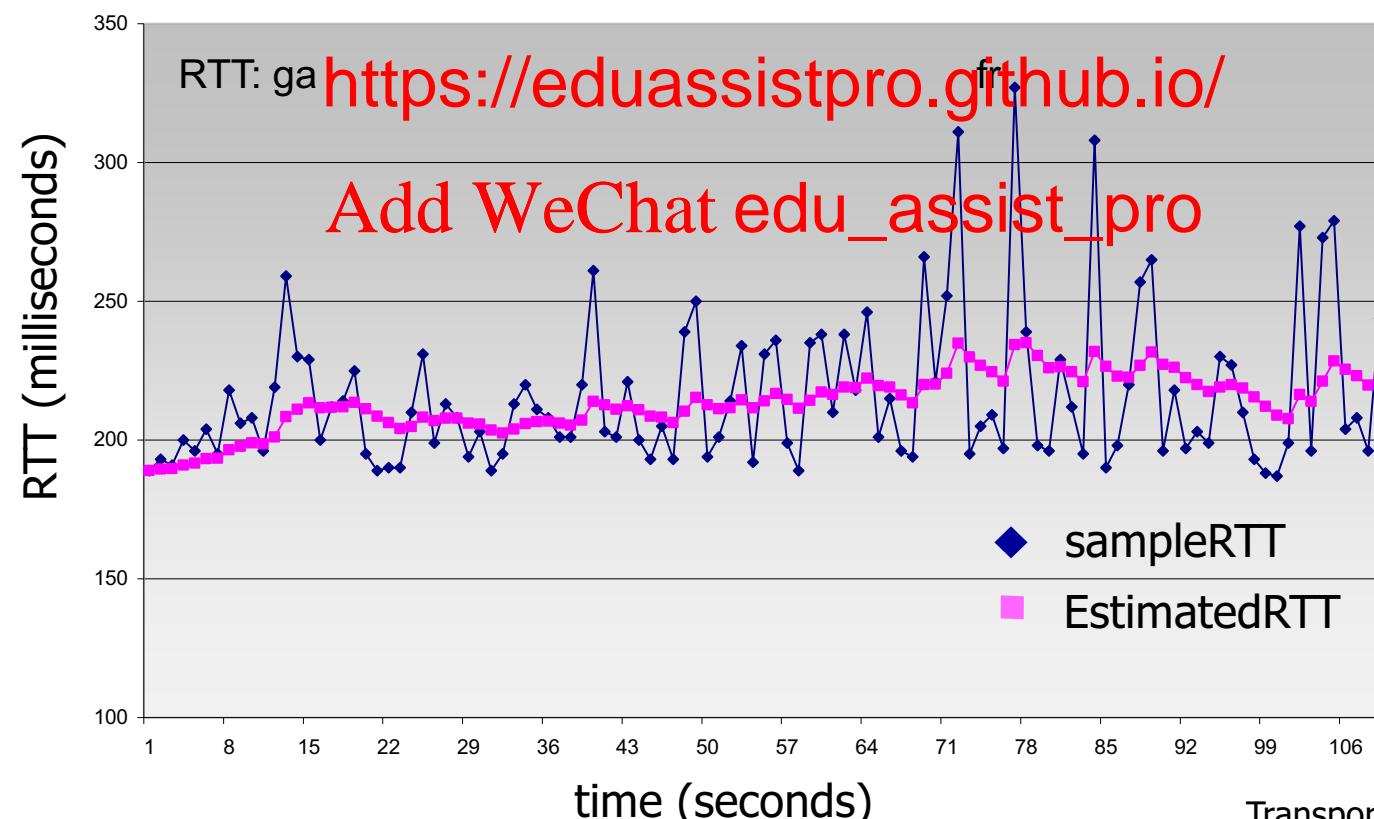
- average several recent measurements, not just current **SampleRTT**

# TCP round trip time, timeout

$$\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$

- ❖ exponential weighted moving average
- ❖ influence of past sample decreases exponentially fast
- ❖ typical value:  $\alpha = 0.125$

Assignment Project Exam Help



# TCP round trip time, timeout

- ❖ **timeout interval:** **EstimatedRTT** plus “safety margin”
  - large variation in **EstimatedRTT** -> larger safety margin
- ❖ estimate SampleRTT deviation from EstimatedRTT:  
**Assignment Project Exam Help**

**DevRTT = (1**

<https://eduassistpro.github.io/>

(typically,  
**Add WeChat edu\_assist\_pro**

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$



↑  
estimated RTT

↑  
“safety margin”

Practice Problem:

[http://wps.pearsoned.com/ecs\\_kurose\\_comnnetw\\_6/216/55463/14198700.cw/index.html](http://wps.pearsoned.com/ecs_kurose_comnnetw_6/216/55463/14198700.cw/index.html)

# TCP round trip time, timeout

(EstimatedRTT+4\*DevRTT)

Assignment Project Exam Help

<https://eduassistpro.github.io/>

RTT  
EstimatedRTT

Add WeChat edu\_assist\_pro

DevRTT

**TimeoutInterval = EstimatedRTT + 4\*DevRTT**



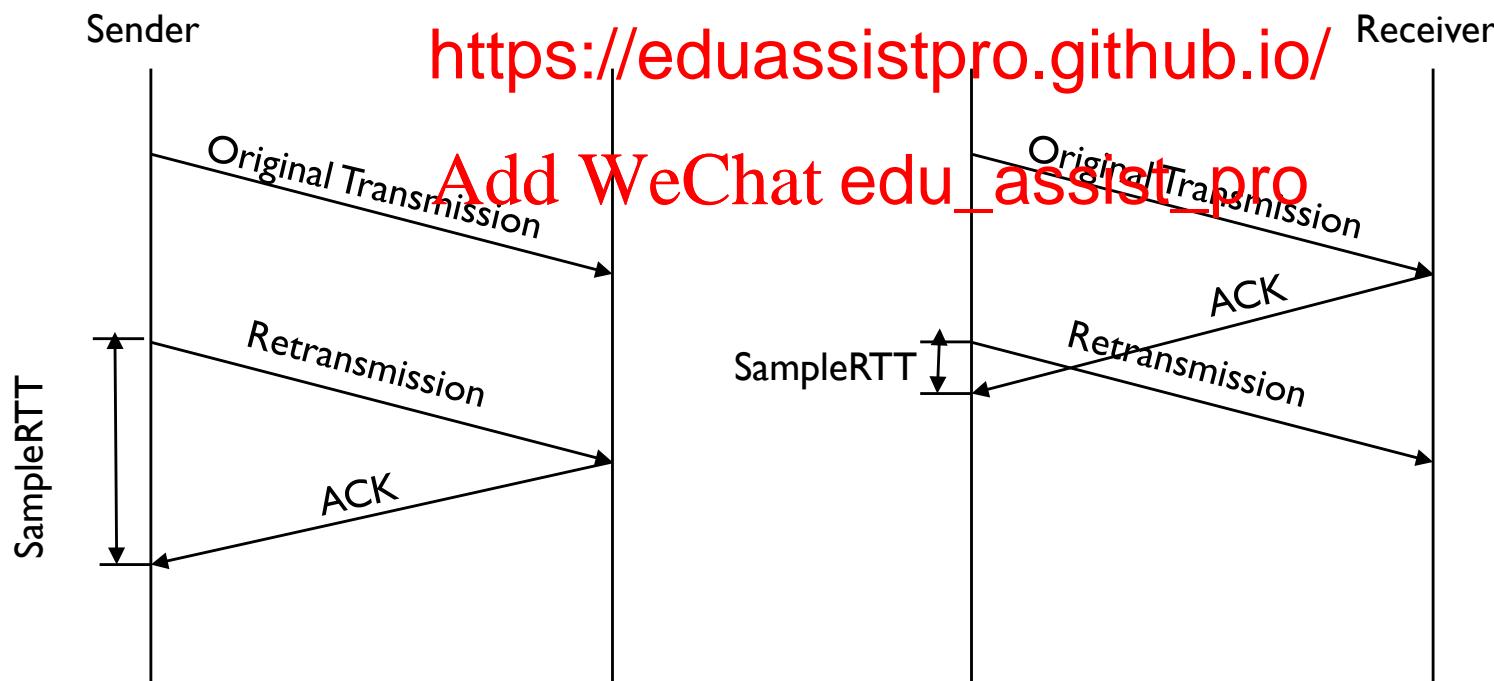
↑  
estimated RTT

↑  
“safety margin”

# Why exclude retransmissions in RTT computation?

- ❖ How do we differentiate between the real ACK, and ACK of the retransmitted packet?

Assignment Project Exam Help



# TCP sender events:

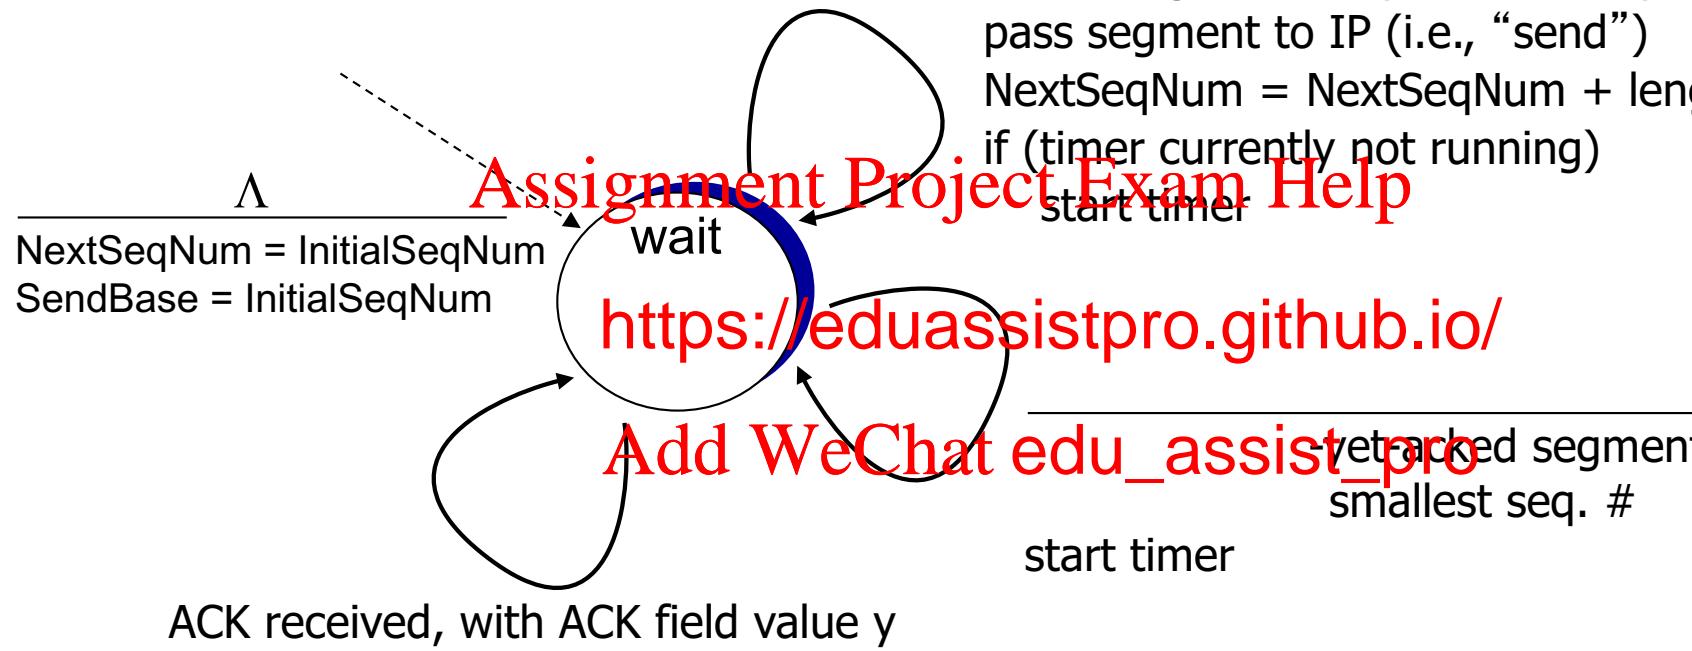
PUTTING IT  
TOGETHER

*data rcvd from app:*

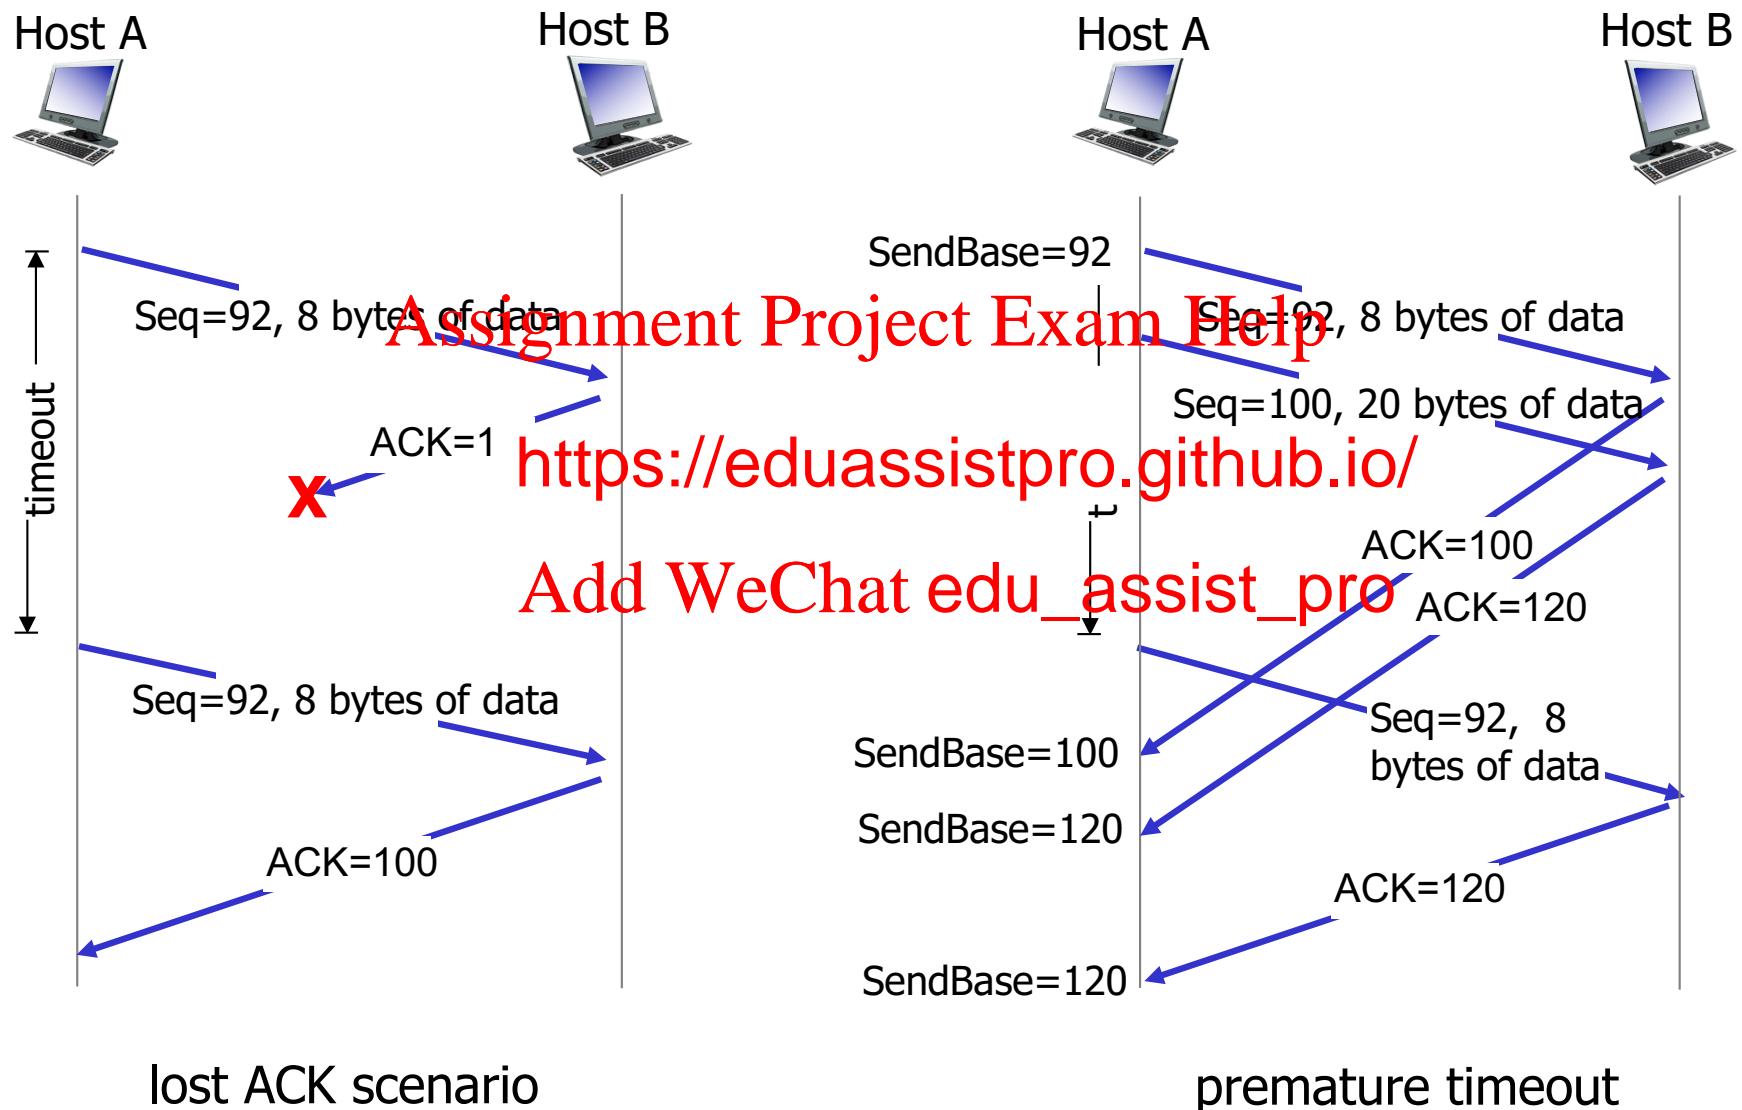
- ❖ create segment with seq #
  - ❖ seq # is byte-stream number of first byte in segm [https://eduassistpro.github.io/  
Assignment Project Exam Help  
knowle](https://eduassistpro.github.io/knowle)
  - ❖ start timer if not already running
    - think of timer as for oldest unacked segment
    - expiration interval: `TimeOutInterval`
- timeout:*
- ❖ retransmit segment that caused timeout
  - ❖ restart timer
  - update what is known to be ACKed
  - start timer if there are still unacked segments

# TCP sender (simplified)

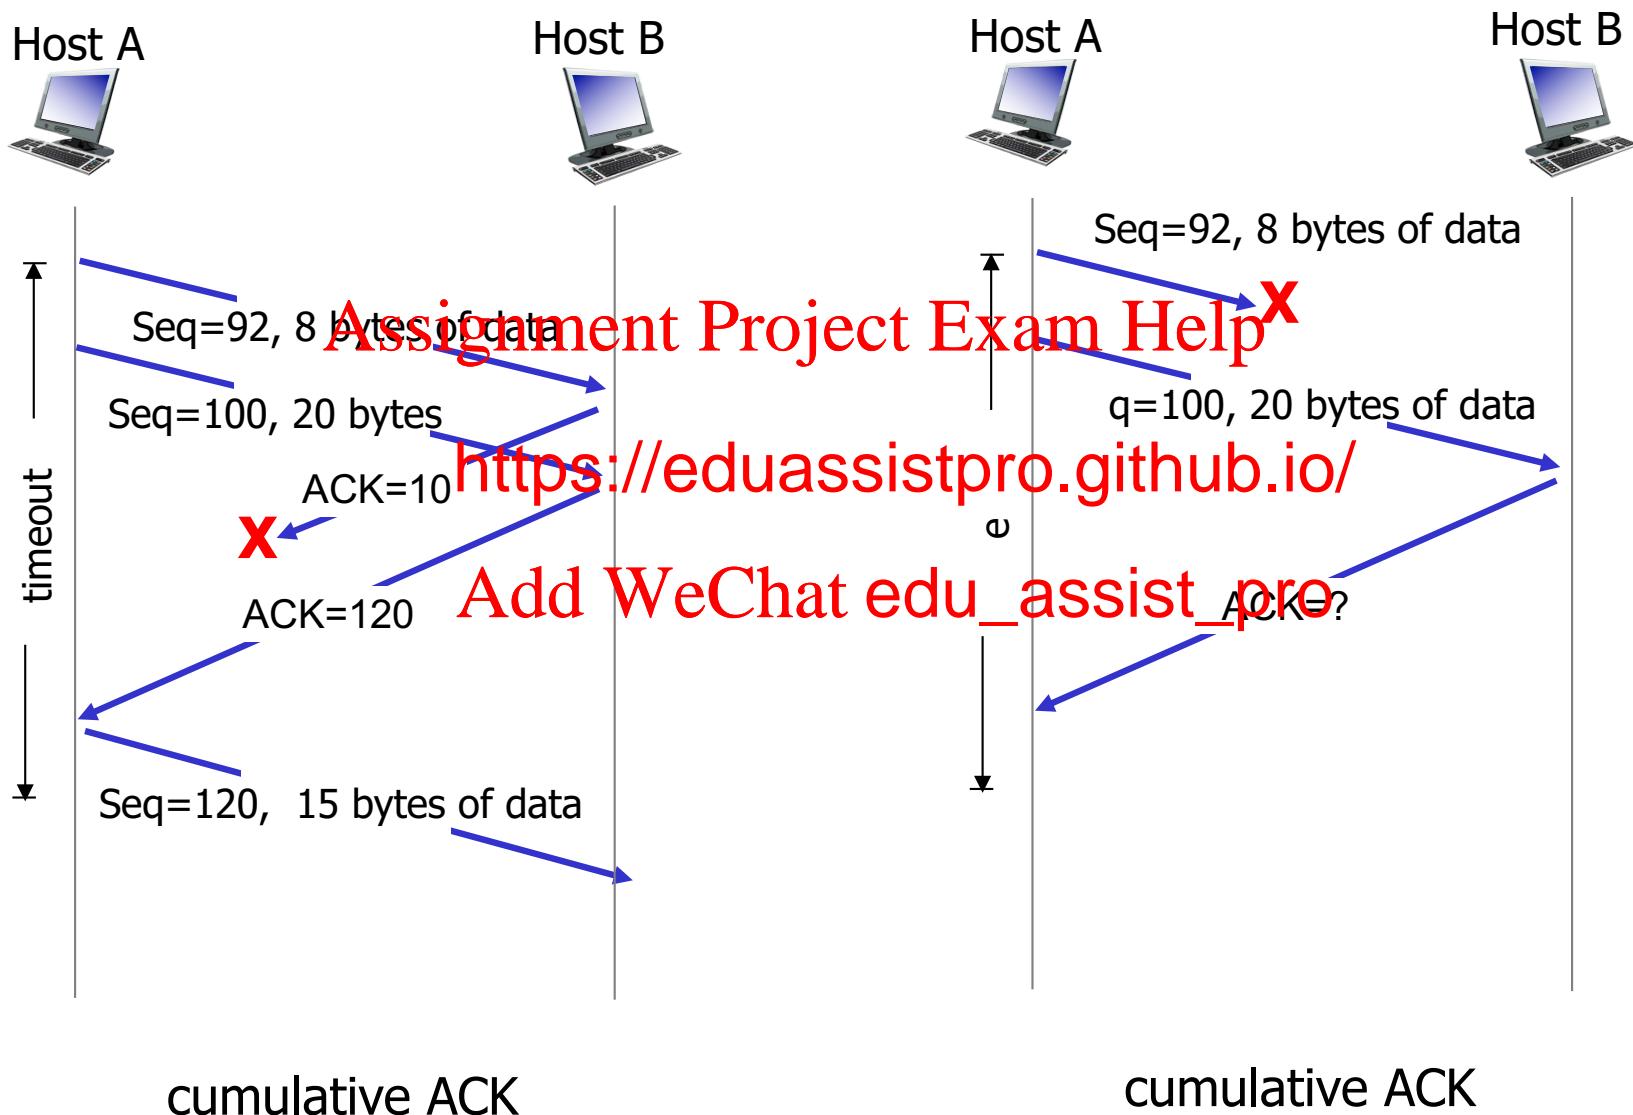
PUTTING IT  
TOGETHER



# TCP: retransmission scenarios



# TCP: retransmission scenarios



# Quiz



Suppose host A sends two TCP segments back to back to Host B over a TCP connection. The first segment has sequence number 90; the second has sequence number 110.

**Assignment Project Exam Help**

- ❖ Q1: How much da <https://eduassistpro.github.io/>?  
**Add WeChat edu\_assist\_pro**
- ❖ Q2: Suppose that the first segment is lost but the second segment arrives at B. In the acknowledgement that Host B sends to Host A, what will be the acknowledgement number?

# TCP ACK generation [RFC 1122, RFC 2581]

<i>event at receiver</i>	<i>TCP receiver action</i>
arrival of in-order segment with expected seq #. All data up to expected seq # already ACKed	delayed ACK. Wait up to 500ms for next segment. If no next segment, send ACK
arrival of in-order seg expected seq #. One other segment has ACK pending	<a href="https://eduassistpro.github.io/">https://eduassistpro.github.io/</a> end single cumulative A $n^{th}$ in-order segments Add WeChat edu_assist_pro
arrival of out-of-order segment higher-than-expect seq. # . Gap detected	immediately send <i>duplicate ACK</i> , indicating seq. # of next expected byte
arrival of segment that partially or completely fills gap	immediate send ACK, provided that segment starts at lower end of gap

# What does TCP do?

Most of our previous tricks, but a few differences

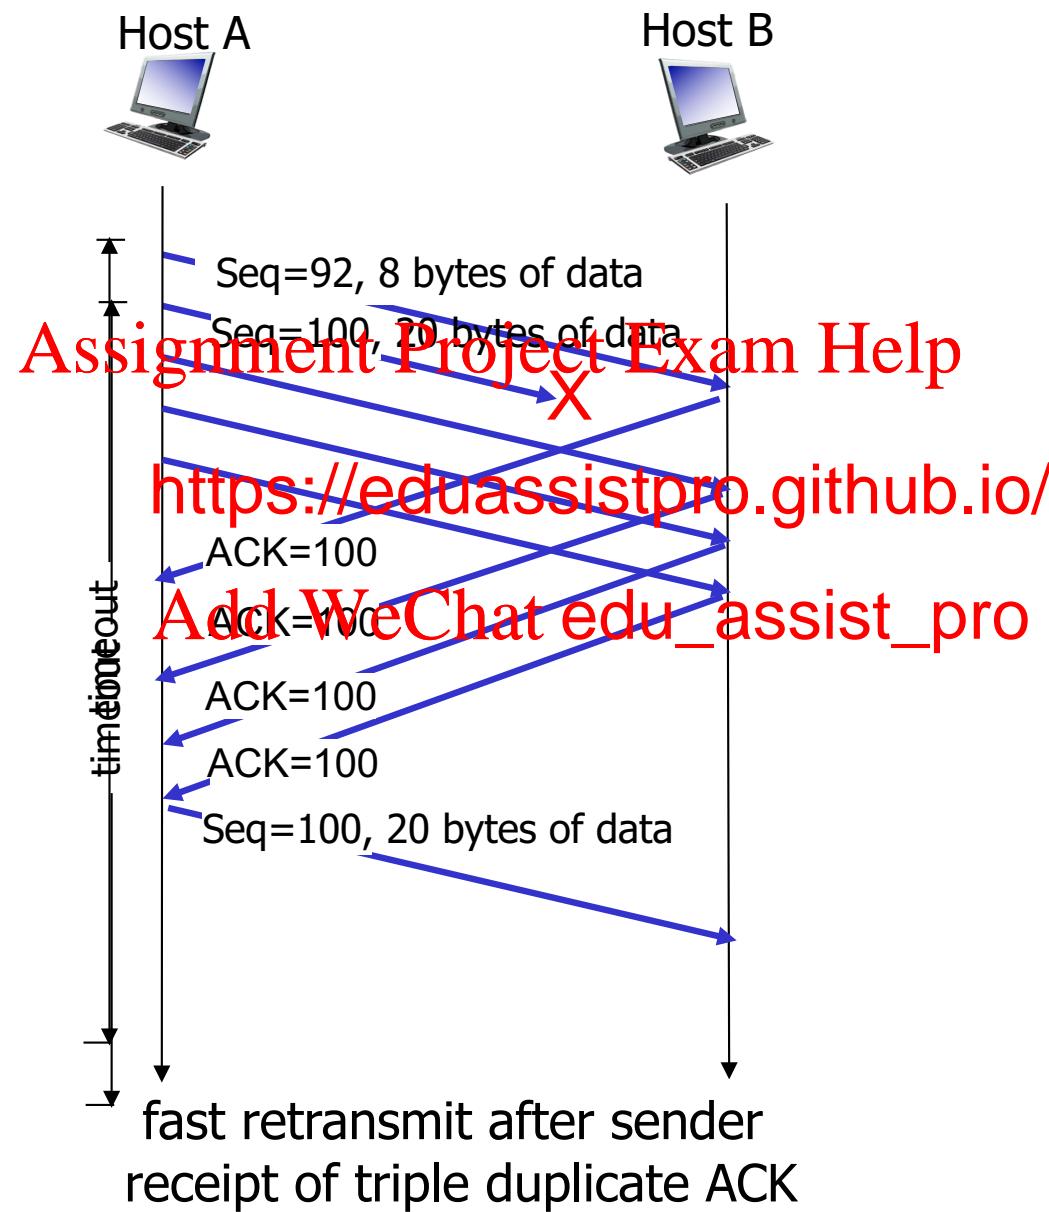
- ❖ Checksum
- ❖ Sequence numbers are byte offsets
- ❖ Receiver sends cumulative acknowledgments (like GBN)
- ❖ Receivers may not acknowledge packets (like SACK)
- ❖ Sender maintains a window (like GBN) and retransmits on timeout
- ❖ Introduces **fast retransmit**: optimisation that uses duplicate ACKs to trigger early retransmission

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat [edu\\_assist\\_pro](https://edu_assist_pro) (like GBN) and

# TCP fast retransmit



# TCP fast retransmit

- ❖ time-out period often relatively long:
  - long delay before resending lost packet
- ❖ “Duplicate ACKs” are a sign of an isolate
  - The lack of ACK means that pack been delivered
  - Stream of ACKs means some packets are being delivered
  - Could trigger resend on receiving “k” duplicate ACKs (TCP uses  $k = 3$ )

## *TCP fast retransmit*

if sender receives 3 duplicate ACKs for a segment with smallest seq #

- likely that unacked segment is lost, so don't wait for timeout

# What does TCP do?

Most of our previous ideas, but some key differences

- ❖ Checksum Assignment Project Exam Help
- ❖ Sequence numbers <https://eduassistpro.github.io/>
- ❖ Receiver sends cumulative acknowledgments (like GBN)
- ❖ Receivers do not drop out-of-sequence packets (like SR)
- ❖ Sender maintains a single retransmission timer (like GBN) and retransmits on timeout
- ❖ Introduces fast retransmit: optimization that uses duplicate ACKs to trigger early retransmission

# Transport Layer Outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

■ segment structure  
able data transfer

https://eduassistpro.github.io/  
Add WeChat edu\_assist\_pro  
connection management  
3.6 es of congestion control

3.7 TCP congestion control

# TCP flow control

application may  
remove data from  
TCP socket buffers ....

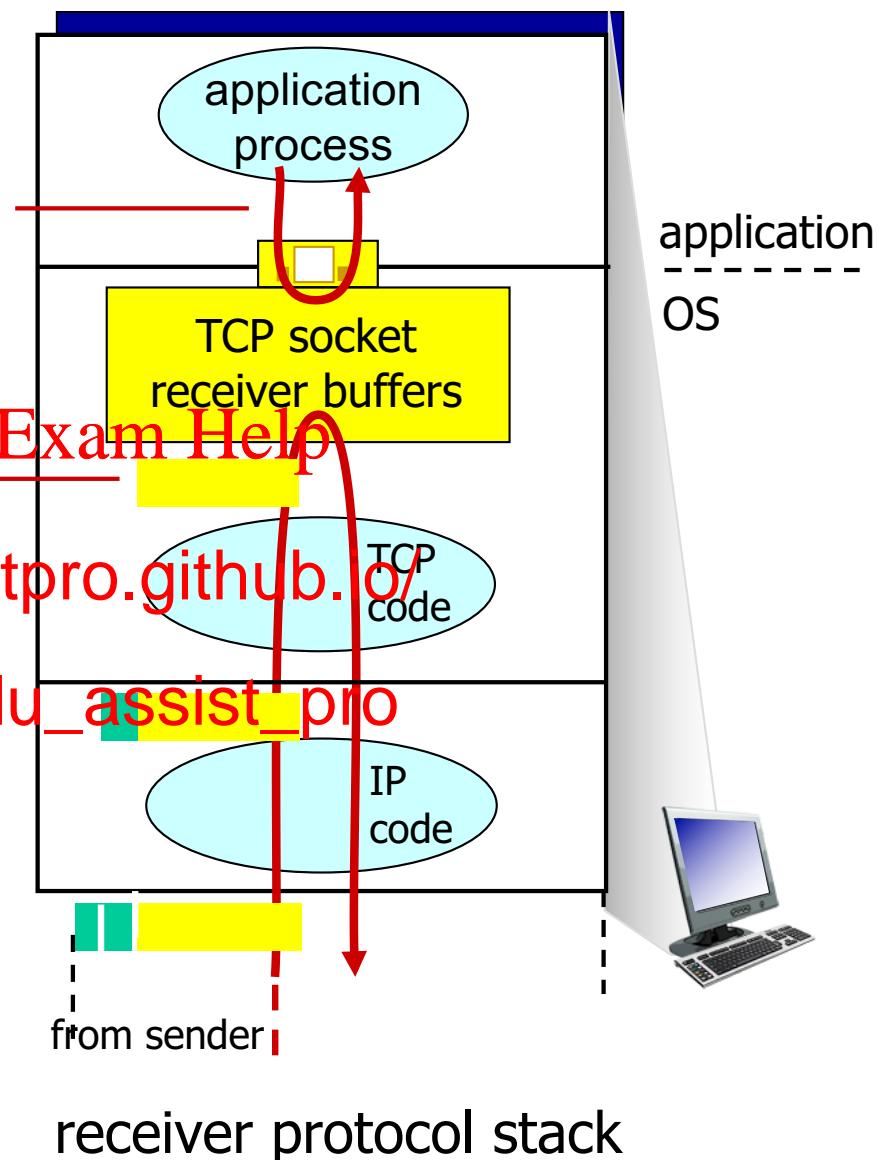
Assignment Project Exam Help  
... slower than TCP receiver is delivering

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

## **flow control**

receiver controls sender, so  
sender won't overflow  
receiver's buffer by transmitting  
too much, too fast

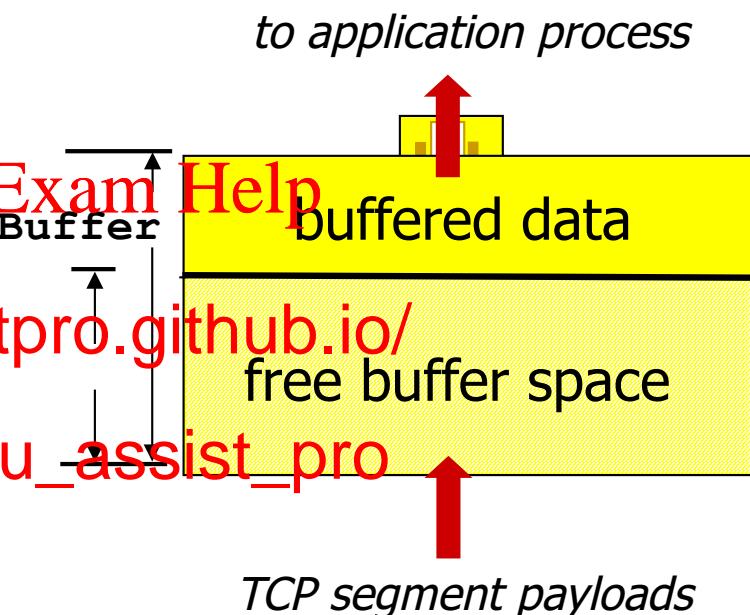


# TCP flow control

- ❖ receiver “advertises” free buffer space by including **rwnd** value in TCP header of receiver-to-sender segments
  - **RcvBuffer** si <https://eduassistpro.github.io/> socket options (is 4096 bytes)
  - many operating systems autoadjust **RcvBuffer**

Assignment Project Exam Help

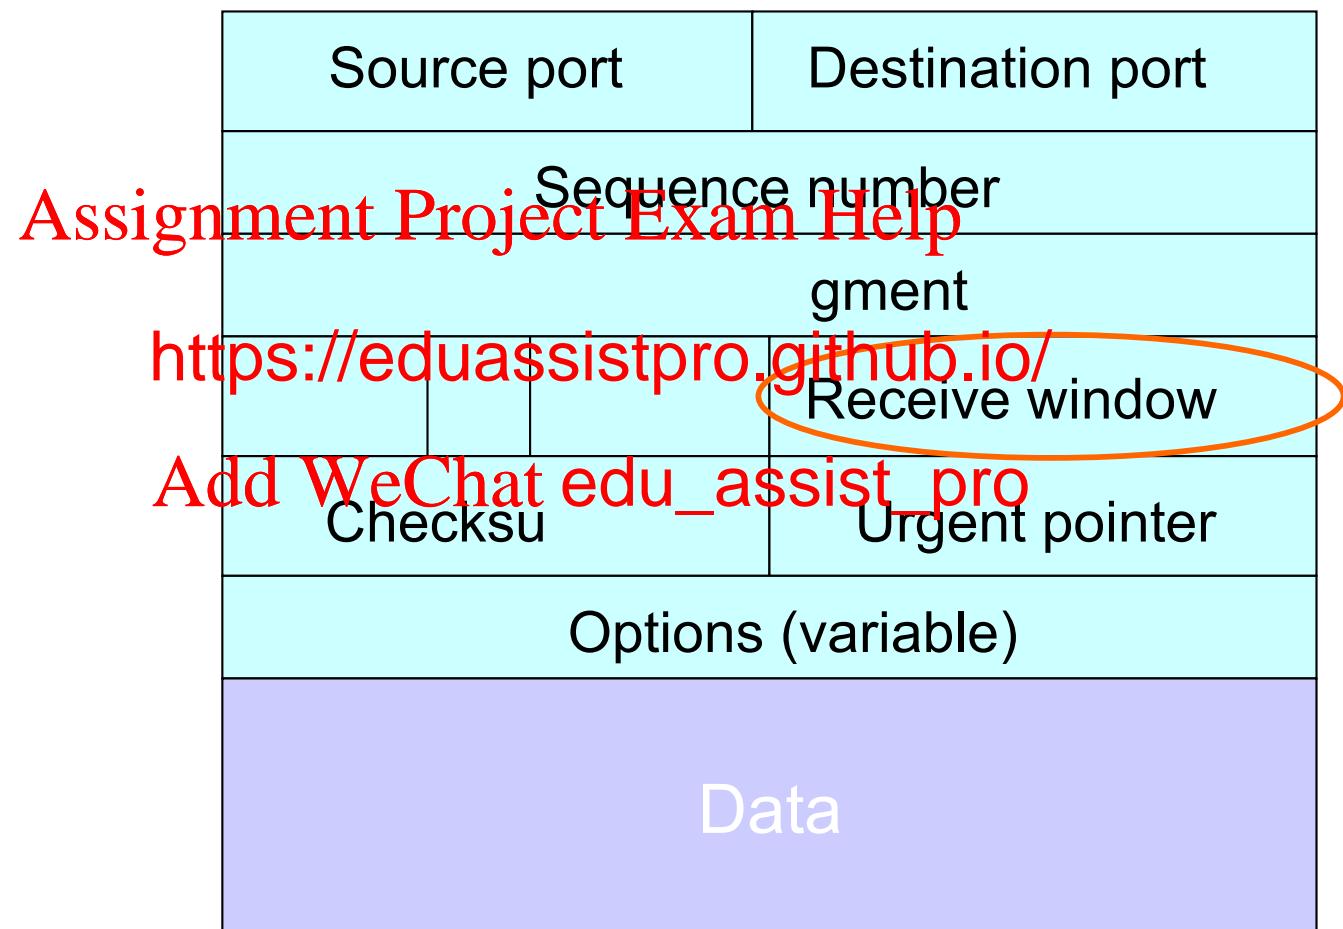
Add WeChat edu\_assist\_pro



- ❖ sender limits amount of unacked (“in-flight”) data to receiver’s **rwnd** value
- ❖ guarantees receive buffer will not overflow

*receiver-side buffering*

# TCP Header



# Transport Layer Outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

■ segment structure  
able data transfer

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

ction management  
es of congestion  
control

3.7 TCP congestion control